

JavaScript BigFloat Library
(Arbitrary precision)
Version 3
By ***Henrik Vestermark (hve@hvks.com)***

Contents

History and Motivation for BigFloat.js.....	7
The First Generation: Decimal String Approach (2012-2015)	7
The Turning Point: BigInt and Key Insights (Fall 2025).....	7
IEEE 754 Inspiration and Binary Representation.....	8
BigFloat Internal Format and Design Philosophy	8
Understanding the Core Structure.....	8
The Philosophy Behind BigInt	9
Why Binary Exponents Make Sense.....	10
The Concept of the Hidden Binary Point.....	10
The Normalization Invariant.....	10
Handling Zero as a Special Case	11
Putting It All Together: The Complete Representation	11
The Decision to Separate the Sign.....	11
The Guarantee of Normalization	12
Dynamic Precision: A Powerful Feature	12
Why Not ES6 Classes?	12
Performance Through Careful Design	13
The Evolution of Bit Length Tracking.....	13
Binary Exponents and Scaling Efficiency	14
Leveraging Native BigInt Performance	14
The Compound Effect.....	15
Understanding Precision Semantics.....	15
The Standard Pattern for Implementing Operations	16
How BigFloat Relates to IEEE 754	16
Developing BigFloat with AI Assistance: A Collaborative Journey	17
The Initial Vision and Communication Challenge	17
The Conceptual Hurdle	17
The Breakthrough Moment.....	18
Rapid Development After Understanding	18
The Iterative Refinement Process	19
Lessons About AI-Assisted Development.....	19
The Value of Domain Knowledge	20
Reflections on the Process	20
BigFloat.....	22
Constructor.....	22
BigFloat Internals.....	23
Object Members.....	23
Normalized Representation.....	23
BigInt for Performance	23
Binary Exponent System.....	23
Implicit Binary Point.....	24
Separate Sign Handling.....	24
Operation Guarantees.....	24

Flexible Precision and Rounding	24
Constructor Design Philosophy	24
Rounding modes	25
Precision.....	25
BigFloat.js Methods and Functions Reference	26
Constructor.....	26
BigFloat(value, precision, rounding)	26
Configuration	26
BigFloat.defaultrounding(mode)	26
BigFloat.defaultprecision(prec)	26
Instance Property Accessors	26
Type Checking Methods	26
Conversion and Formatting Methods.....	26
Basic Arithmetic Operations.....	27
Comparison Operations	27
Rounding and Integer Functions.....	27
Special Mathematical Functions	27
Adjacent Value Functions.....	28
Mathematical Constants.....	28
Power and Exponential Functions	28
Trigonometric Functions.....	28
Hyperbolic Functions.....	28
Utility and Internal Methods.....	29
Rounding Modes.....	29
Usage Examples	29
Basic arithmetic:	29
Computing π to 100 digits:	29
Trigonometric calculations:	29
Using fma for accurate computation:.....	29
Using nextafter:.....	30
BigFloat.js API Reference	30
Overview	30
Constants and Configuration.....	30
Mathematical Constants.....	30
Instance Properties	30
Special Value Constants	31
Rounding Mode Constants.....	31
Default Settings.....	31
Methods and Functions	32
BigFloat.abs().....	32
BigFloat.acos().....	32
BigFloat.acosh().....	32
BigFloat.add()	32
BigFloat.asin()	34
BigFloat.asinh()	34
BigFloat.assign()	34

BigFloat.atan()	34
BigFloat.atan2()	35
BigFloat.atanh()	35
BigFloat.ceil()	35
BigFloat.clone()	35
BigFloat.cos()	36
BigFloat.cosh()	36
BigFloat.defaultPrecision()	36
BigFloat.defaultRounding()	36
BigFloat.div()	37
BigFloat.E()	37
BigFloat.equal()	37
BigFloat.EPSILON()	37
BigFloat.exp()	39
BigFloat.floor()	39
BigFloat.fma()	39
BigFloat.fmod()	39
BigFloat.greater()	40
BigFloat.greaterequal()	40
BigFloat.inverse()	40
BigFloat.isFinite()	41
BigFloat.isInfinite()	41
BigFloat.isInteger()	41
BigFloat.isNaN()	42
BigFloat.isNegative()	42
BigFloat.isOdd()	42
BigFloat.isPositive()	42
BigFloat.ldexp()	43
BigFloat.less()	43
BigFloat.lessequal()	43
BigFloat.LN10()	43
BigFloat.log()	44
BigFloat.log10()	44
BigFloat.mul()	46
BigFloat.neg()	46
BigFloat.nextafter()	46
BigFloat.notequal()	46
BigFloat.pow()	47
BigFloat.precision()	47
BigFloat.pred()	47
BigFloat.rounding()	48
BigFloat.sign()	48
BigFloat.sin()	48
BigFloat.sqrt()	49
BigFloat.SQRT2()	49
BigFloat.sub()	49

BigFloat.succ()	49
BigFloat.tanh()	50
BigFloat.toBigInt()	50
BigFloat.toExponential()	50
BigFloat.toNumber()	51
BigFloat.toPrecision()	51
BigFloat.toString()	51
parseBigFloat()	52
BigFloat.toInteger()	52
BigFloat.toNumber()	52
BigFloat.toString()	53
BigFloat.trunc()	53
Appendix	54
Part 1: BigFloat vs BigNumber.js	54
Internal Representation	54
BigNumber.js Structure	54
BigFloat Structure	55
Performance Implications	55
Arithmetic Operations	55
Conversion Overhead	56
Precision Handling	56
BigNumber.js Approach	56
BigFloat Approach	56
Mathematical Function Coverage	57
BigNumber.js Functions	57
BigFloat Functions	57
API Design and Usability	57
BigNumber.js API	58
BigFloat API	58
Memory and Performance Characteristics	58
Memory Footprint	58
Speed Comparisons	59
Error Handling and Edge Cases	59
BigNumber.js Approach	59
BigFloat Approach	59
Part 2: BigFloat vs Decimal.js	60
Internal Representation	60
Decimal.js Structure	60
Mathematical Function Coverage	60
Decimal.js Functions	60
Implementation Approaches	61
Mathematical Constants: A Critical Difference	61
Specialized Numerical Functions	61
Performance Characteristics	63
Arithmetic Speed	63
Transcendental Functions	63

Algorithm Modernity and Optimization	64
Precision and Accuracy.....	65
Decimal.js Precision Model.....	65
BigFloat Precision Model	65
API and Usability.....	66
Decimal.js API.....	66
BigFloat API Comparison.....	66
Maturity and Ecosystem	66
Decimal.js Maturity	66
BigFloat's Position	67
Code Size and Complexity.....	67
Conclusion: Choosing the Right Library	68
Choose BigNumber.js When:	68
Choose Decimal.js When:.....	68
Choose BigFloat When:.....	68
The Fundamental Trade-offs.....	69
Maturity Considerations.....	69
Looking Forward	69
Part 3. Performance Comparison Analysis	71
Executive Summary	71
1. Square Root: $\sqrt{2}$).....	71
2. Natural Logarithm: $\ln(2.5)$	71
3. Base-10 Logarithm: $\log_{10}(2.5)$	72
4. Exponential Function: $e^{2.5}$	72
5. Sine Function: $\sin(1.5)$	72
6. Tangent Function: $\tan(1.5)$	73
7. Inverse Sine: $\arcsin(0.5)$	73
8. Hyperbolic Sine: $\sinh(1.5)$	73
9. Hyperbolic Tangent: $\tanh(1.5)$	74
10. Inverse Hyperbolic Sine: $\text{asinh}(0.5)$	74
Comprehensive Analysis	75
Performance Scaling Characteristics	75
Reliability and Robustness.....	75
Most Significant Performance Advantages	75
Areas of Competitive Performance.....	76
Conclusions and Recommendations	77
Overall Assessment.....	77
Technical Advantages.....	77
Recommended Use Cases	77
When Decimal.js Might Be Considered	78
Final Recommendation	78
Technical Notes	80
Benchmark Methodology.....	80
Error Conditions.....	80
About This Analysis	80

History and Motivation for BigFloat.js

The First Generation: Decimal String Approach (2012-2015)

I have maintained a mature C++ arbitrary-precision library for many years, and it was time to make arbitrary-precision arithmetic available in JavaScript. I wanted it to be simple and easy to debug. At the time, debugging JavaScript was cumbersome, so I decided to store numbers internally as decimal strings. This made sense at that time because I could immediately see the value of any number as a decimal string, allowing me to correct mistakes as I progressed through the project. Keeping the internal representation as decimal strings also avoided any conversion to and from the internal format, greatly simplifying the code, though at the expense of severe performance degradation compared to using a binary representation closer to the underlying hardware architecture.

It worked. I quickly got a functional JavaScript library. However, performance-wise, it fell short of BigNumber.js and Decimal.js, as expected. I examined BigNumber.js and noticed that it stores 14 decimal digits in a single JavaScript float variable, making its representation considerably more compact and much faster. My original library's strengths were that it included auxiliary trigonometric and hyperbolic functions, plus it could generate arbitrary precision constants for π , $\ln(2)$, $\ln(10)$, and e , more than what BigNumber supported, and far beyond Decimal.js's artificial limitation of these constants to only 1,025 decimal digits. However, the lack of performance made it more of a personal internal project.

The Turning Point: BigInt and Key Insights (Fall 2025)

After many years of pondering a complete rewrite, the time came in the fall of 2025 to start fresh. I felt particularly compelled after JavaScript added support for arbitrary-precision integers via the BigInt type. I knew from experience that considerable performance gains could be achieved since most of the heavy lifting for the three basic operations ($+$, $-$, \times) could be performed using BigInt and therefore executed internally inside the JavaScript interpreter (which, to my knowledge, is written in C++), instead of through a series of interpreted JavaScript statements.

Another key insight was that integer and floating-point multiplication can use the same underlying multiplication algorithm, the only difference is that floating-point has a decimal point between the integer and fractional parts. For example, if you multiply 123×456 , you get 56,088. If you use the same digits in a floating-point multiplication, say 1.23×45.6 , you get 56.088. The digits are still the same and in the same order; with floating-point multiplication, you have to determine where the decimal point should be placed within the resulting digits.

Notice that division is not critical here. In floating-point arithmetic, you typically turn to Newton or Goldschmidt iteration methods that only require the availability of $+$, $-$, and \times operations. If an efficient implementation using the BigInt JavaScript type could be achieved, then all the remaining pieces, like exponential, logarithm, trigonometric, and hyperbolic functions, could be built using these three underlying building blocks.

IEEE 754 Inspiration and Binary Representation

By choosing BigInt as the internal type to store arbitrary-precision floating-point numbers, I also had to accept a binary (base 2) internal representation. This design choice draws inspiration from the IEEE 754 floating-point standard, which has proven remarkably successful in hardware implementations. Like IEEE 754, BigFloat represents numbers as a sign, a binary significand (mantissa), and a binary exponent. The number's value is computed as: sign \times significand \times 2^{exponent} .

This binary representation aligns naturally with how computers operate at the hardware level and enables BigFloat to leverage the highly optimized BigInt operations provided by JavaScript engines. The significand is normalized to have its leading bit set, representing a value in the range [1, 2), similar to how IEEE 754 normalizes its mantissa.

The binary exponent (representing powers of 2) makes scaling operations trivial.

Multiplying by powers of 2 is simply exponent addition, and detecting powers of 2 is immediate.

The trade-off, of course, is that this binary representation requires conversion when interfacing with the decimal human world. Converting decimal numbers to their internal binary representation for operations, then converting back to decimal for display, adds computational overhead. However, this cost is more than offset by the dramatic performance gains in arithmetic operations, especially for intensive numerical computations, where many operations are performed on each number.

With these insights in place, leveraging BigInt for performance, using an IEEE 754-inspired binary representation for efficiency, and accepting the overhead of decimal-binary conversion, I could now start my project to rewrite my BigFloat JavaScript library from scratch.

BigFloat Internal Format and Design Philosophy

Understanding the Core Structure

At its heart, a BigFloat object is composed of several carefully chosen members that work together to represent arbitrary-precision floating-point numbers. Each member plays a specific role in maintaining both the mathematical correctness and computational efficiency of the representation.

The **_sign** member stores whether the number is positive or negative, using the convention +1 for positive and -1 for negative. By keeping the sign separate from the numerical content, we make many operations more straightforward and the code easier to understand.

The **_exponent** represents the power of 2 by which the significand is scaled. To determine the actual contribution of this exponent to the final value, we compute 2^{exponent} . This binary exponent approach was deliberately chosen to align with the principles of the IEEE 754 floating-point standard and provides highly efficient scaling

operations, since adjusting a number's magnitude often requires nothing more than adding or subtracting from the exponent.

The `_precision` field indicates how many decimal fraction digits we maintain for a normalized number. When we talk about a normalized number, we're referring to a representation in the format 1.xxxxx, which means the significand always represents a value in the range from 1 up to (but not including) 2. This normalization convention ensures that every number has a consistent, canonical representation and that we use our available precision efficiently.

At the core of every BigFloat sits the `_significand`, which holds the actual numerical content of our number. This is stored as a JavaScript BigInt object, giving us access to arbitrary-precision integer arithmetic without implementing it ourselves. The significand is where the real "digits" of our number reside, encoded in binary.

The `_rounding` member specifies the rounding method to use when normalizing a number after an operation. Different mathematical and scientific applications require different rounding behaviors, so we support several standard modes including "nearest" (which rounds to the nearest value, with ties going to the even number), "down" (which rounds toward zero), "up" (which rounds away from zero), "floor" (which rounds toward negative infinity), and "ceiling" (which rounds toward positive infinity).

The `_special` flag handles those exceptional values that don't fit into the standard numerical representation: zero, NaN (Not a Number), positive infinity, and negative infinity. When this flag is set to indicate one of these special values, the significand and exponent may not follow their usual interpretation rules.

Finally, the `_bits` field serves as a performance optimization by caching the number of bits in the significand. Counting bits can be expensive if we need to do it repeatedly, so we cache the result. Whenever the `_bits` field is negative or undefined, we know we need to recalculate it by actually counting the bits in the significand.

The Philosophy Behind BigInt

One of the most consequential design decisions in BigFloat is the choice to use JavaScript's native BigInt object for storing the significand. This decision dramatically improves performance and eliminates the enormous complexity that would otherwise plague the implementation. Without BigInt, we would need to manage arrays or vectors of smaller numbers to hold our arbitrary-precision values, constantly juggling carries, borrows, and the bookkeeping involved in multi-precision arithmetic.

By leveraging BigInt, we gain access to native arbitrary-precision integer arithmetic that's been carefully optimized by JavaScript engine developers. We get efficient bitwise operations essentially for free. Memory management is handled automatically, so we don't need to allocate or deallocate storage for digit arrays. Perhaps most importantly,

BigInt integrates directly with JavaScript's type system, making it feel natural to work with rather than a bolt-on add-on.

Why Binary Exponents Make Sense

The decision to handle exponents in base 2 rather than base 10 makes the internal structure remarkably similar to the IEEE 754 floating-point format that underlies most modern computer arithmetic. While BigFloat isn't identical to IEEE 754, it offers arbitrary precision; this design choice provides numerous practical advantages.

Multiplication and division become much simpler operations because exponents can be added or subtracted. If you multiply two numbers with exponents e_1 and e_2 , the result has exponent $e_1 + e_2$, regardless of what the significands are doing. Powers of 2 are trivially easy to detect and work with. The normalization logic is straightforward because shifting the significand left or right by one bit corresponds exactly to incrementing or decrementing the exponent by one. All of these properties align naturally with how computers actually work at the hardware level.

The Concept of the Hidden Binary Point

Understanding how BigFloat interprets the significand requires an understanding of the hidden binary point. We don't treat the `_significand` as just a plain BigInt number representing an integer value. Instead, we conceptualize it as an implicit binary point, like a decimal point, positioned between the first set bit and the remaining bits.

Let me illustrate this with a concrete example. Suppose the `_significand` contains the value `0b1101`, which is 13 in decimal. We don't interpret this as the integer thirteen. Instead, we imagine a binary point after the first bit: 1.101 in binary notation. If we convert this to decimal, $1.101_{\text{binary}} = 1 + 0.5 + 0.125 = 1.625$.

This interpretation is absolutely crucial to understanding how arithmetic operations work in BigFloat. When we multiply, add, or subtract BigFloat numbers, the actual bit patterns that emerge from these operations on the significands are mathematically correct. The key insight is that we only need to track where the hidden binary point ends up after the operation completes, and then perform a normalization step to ensure that there's always exactly one bit set before (to the left of) that binary point.

The Normalization Invariant

Every BigFloat number, except for special cases like zero, maintains a critical invariant: the significand, when interpreted with its hidden binary point, must represent a value in the range from 1 up to (but not including) 2. This means the most significant bit of the significand is always 1, except in the special case of zero.

Let's look at how various numbers appear in this normalized form. The number 5.0 is represented as 1.01 in binary multiplied by 2 squared, so we store a significand of `0b101`

with an exponent of 2. The number 0.75 becomes 1.1 in binary multiplied by 2 to the power of -1, giving us a significand of 0b11 and an exponent of -1. For 12.5, we have 1.1001 in binary, multiplied by 2^3 , resulting in a significand of 0b11001 and an exponent of 3.

This normalization ensures that every number has a unique, canonical representation. There's only one way to represent any given value (aside from differences in precision and rounding), which makes comparisons and other operations much more reliable and efficient.

Handling Zero as a Special Case

Zero presents an interesting challenge because it fundamentally cannot be represented in the normalized [1, 2) format. No matter how you try to write zero, you cannot express it as something in the range [1, 2) multiplied by any power of 2. Therefore, zero is treated as a special case, indicated by setting the `_significand` to 0n (the BigInt value zero) and marking it with the appropriate `_special` flag.

Putting It All Together: The Complete Representation

To summarize how we represent numbers in BigFloat, we use three primary pieces of information working in concert. The `_significand` contains the binary mantissa, with the implicit binary point positioned after the first bit. The exponent indicates the base-2 scaling factor; we multiply by 2 raised to this power. The `_sign` indicates whether the overall value is positive or negative.

When we want to compute the actual numerical value that a BigFloat represents, we calculate: `sign × (significand interpreted as 1.xxxxx in binary) × 2exponent`. This formula captures the complete representation scheme.

The Decision to Separate the Sign

You might wonder why we keep the sign separate from the significand. After all, JavaScript's BigInt fully supports negative numbers using two's complement representation, so we could theoretically store negative significands directly. However, I deliberately chose to handle the sign separately for several compelling reasons.

Keeping the sign separate makes comparison logic much more straightforward. When we want to compute the absolute value of a number, we need to set the sign to +1 rather than manipulating the significand. Operations that depend on the sign become more explicit and easier to understand. From a debugging and inspection perspective, it's much clearer to see a separate sign field than to have to decode whether a significand is negative. Overall, the code becomes more maintainable and less prone to subtle bugs.

The Guarantee of Normalization

One of the most important invariants in the BigFloat system is that every arithmetic or mathematical function guarantees that its result is returned as a properly normalized number. This normalization encompasses several aspects working together.

The exponent is adjusted so that the significand falls into the [1, 2) range. The sign is set correctly to reflect whether the result is positive or negative. The special flag is updated if the result is zero, NaN, or infinity. Most importantly, the result is rounded to the specified precision using the specified rounding mode, ensuring we don't carry more bits than required.

This invariant dramatically simplifies operation implementation because you can always assume your inputs are properly normalized. As long as you normalize your outputs before returning them, the system remains consistent and predictable.

Dynamic Precision: A Powerful Feature

One of BigFloat's most valuable features is that precision and rounding mode are not fixed when you create a number. You can change them at will on any BigFloat object, giving you tremendous flexibility in how you manage numerical accuracy.

This flexibility becomes extremely valuable when performing complex arithmetic operations. You can increase the precision during intermediate calculations to account for accumulated rounding errors and maintain accuracy where it matters most. After completing the sensitive computations, you can restore the result's precision to your final target level. Different parts of a single calculation can use different precision levels as needed, allowing you to optimize the balance between accuracy and performance.

For example, imagine you're computing the sine of a number, and you want 20 digits of accuracy in your final answer. You might create your initial value with 20 digits of precision, but then increase the precision to 50 digits before calling the sine function to account for the accumulated errors in the Taylor series computation. After getting the high-precision result, you reduce the precision back to 20 digits and normalize, which applies the appropriate rounding to give you your final answer.

Why Not ES6 Classes?

I made a conscious decision not to use JavaScript's ES6 class syntax to implement BigFloat, opting instead for a more straightforward function-based approach. Several practical considerations drove this choice.

Function-based construction offers more flexibility. You can create BigFloat objects using straightforward function calls or object notation, and you can set up factory functions that feel natural to use. The syntax resembles C++ cast operators or type conversion functions, making the code feel familiar to developers from other languages.

This approach also enables very natural-looking expressions. You can write things like `BigFloat.mul(BigFloat(10.5), BigFloat.sqrt(a))`, which reads almost like mathematical notation. This mirrors how JavaScript's built-in `BigInt` works; you write `BigInt(105)` to convert 105 into a `BigInt` object, and `BigFloat` follows the same pattern.

From a practical standpoint, avoiding class syntax makes the code more accessible and easier to integrate with various JavaScript environments and tooling. Function-based APIs are simpler to extend, wrap, and adapt to different use cases without getting tangled up in inheritance complexity. The overall result is code that's easier to understand, modify, and use across different contexts.

Performance Through Careful Design

The design of `BigFloat` incorporates several mechanisms specifically aimed at achieving good performance while maintaining flexibility and correctness.

The Evolution of Bit Length Tracking

One of the most instructive examples of performance optimization in `BigFloat` is the `_bits` field, which caches the bit length of the significand. While this might seem like a minor optimization at first glance, it actually matters quite a bit and illustrates how seemingly small decisions can have significant performance implications.

Initially, when we needed to know how many bits a significand contained, we calculated it on demand every single time using the expression:

```
const bitLength = this._significand.toString(2).length;
```

This approach seems OK at first glance. After all, how expensive can converting a number to a binary string and checking its length really be? The answer, as it turns out, is "surprisingly expensive when you do it thousands or millions of times."

The problem is that converting a `BigInt` to its binary string representation isn't a trivial operation, especially for large numbers with hundreds or thousands of bits. The JavaScript engine must traverse the internal representation of the `BigInt`, convert it to binary digits, allocate a string to store those digits, and populate the string character by character. For a significand with, say, 1000 bits, this means creating a 1000-character string to count its length, then immediately discarding that string.

What made this particularly painful was that we often needed the bit length multiple times during a single operation. Normalization requires the bit length to determine the number of shifts to perform on the significand. Precision adjustments are needed to decide whether to round. Comparison operations use it to quickly estimate relative magnitudes. In a complex calculation involving dozens or hundreds of operations, we might be recalculating the same bit length dozens of times for numbers whose significands hadn't even changed.

The performance impact became obvious during benchmarking. Operations that should have been fast were spending a measurable fraction of their time just counting bits. Profiling revealed that the `toString(2)` calls were appearing prominently in the hot paths of the code.

The solution was elegantly simple: cache the bit length. By adding the `_bits` field to store this value, we transformed an $O(n)$ operation that happened constantly into an $O(1)$ lookup that happened just as often, with the actual calculation only occurring when necessary. Counting bits in an arbitrary-precision integer can be expensive if you have to do it repeatedly, so by caching this value, we avoid that overhead entirely.

The cached value is computed once when the significand is first set or modified, and subsequent accesses read the stored value. When the field is set to a negative value, it signals that we should recalculate it the next time someone needs the bit length; for example, after cloning an object or when we're uncertain whether the cached value is still valid.

This single optimization had a significant impact on performance, particularly for high-precision arithmetic, where the significands are large, and the operations are complex. Benchmarks showed speedups of 20-30% on some common operations, purely from eliminating redundant bit-length calculations. For operations like sine and cosine that perform many internal iterations, the improvement was even more pronounced.

Binary Exponents and Scaling Efficiency

Using base-2 exponents enables efficient multiplication and division. The exponent part of these operations reduces to simple addition or subtraction, while the significand operations proceed independently. This is dramatically faster than manipulating base-10 exponents, which would require conversions and more complex scaling operations.

When you multiply two numbers, their exponents add together. When you divide, the exponents subtract. Need to multiply by a power of two? Just add to the exponent. This operation is essentially free compared to actually manipulating the significand. Even scaling by large factors. E.g., multiplying by 2^{100} , for instance, is just a single integer addition.

The binary exponent system also makes powers of two trivially cheap to detect and work with. If the significand equals 1, you know immediately that you have an exact power of two, and many operations can take specialized fast paths. Square roots of perfect powers of two with even exponents don't even require iteration; you halve the exponent, and you're done.

Leveraging Native BigInt Performance

By leveraging native `BigInt` for arithmetic, we ensure that the actual number crunching occurs in highly optimized engine code rather than in JavaScript loops. The JavaScript

engines V8 in Chrome and Node.js, SpiderMonkey in Firefox, and JavaScriptCore in Safari have invested significant engineering effort to optimize BigInt operations. They use sophisticated algorithms such as Karatsuba multiplication for medium-sized numbers and more advanced techniques, such as Toom-Cook or FFT-based multiplication, for truly large values.

We benefit from this optimization work without implementing it ourselves. When we multiply two BigInt significands, the engine automatically chooses the best algorithm based on the size of the operands. When we perform bitwise operations, the engine uses optimized native code paths. Even memory management for arbitrary-precision values is handled by carefully tuned allocators that minimize fragmentation and maximize cache efficiency.

This represents a huge win compared to implementing arbitrary-precision arithmetic from scratch in JavaScript. Even a well-written pure JavaScript implementation would struggle to match the performance of native BigInt, simply because the engine can use techniques, such as inline assembly, SIMD instructions, and hand-tuned algorithms, that aren't available to JavaScript code.

The Compound Effect

What makes these optimizations particularly effective is how they compound. The `_bits` caching reduces overhead in normalization. Efficient normalization makes lazy normalization more attractive because each deferred normalization saves more work. Binary exponents make the exponent arithmetic within normalization essentially free. Native BigInt performance means the significand operations that dominate the computation time are as fast as we can reasonably make them.

The result is a library that, while not as fast as hardware floating-point for fixed-precision arithmetic, delivers respectable performance even for very high-precision arithmetic. Users can request hundreds or thousands of digits of precision and still get results in a reasonable time. Careful attention to performance details makes the difference between a library that's academically interesting but impractical for real-world use and one that can solve real computational problems requiring extended precision.

Understanding Precision Semantics

An important subtlety in BigFloat is that the `_precision` field represents decimal precision, not binary precision. This distinction matters because users typically think in decimal terms. When someone says they want 50 digits of precision, they mean 50 decimal digits, not 50 binary bits.

The relationship between decimal and binary precision is logarithmic. To represent a given decimal precision in binary, we need approximately $\text{_precision} \times \log_2(10)$ bits, which is about $\text{_precision} \times 3.32$ bits. When we normalize a number, we ensure that the significand has enough bits to represent the requested decimal precision faithfully. The

rounding operation then truncates or adjusts the significand to match the target precision, discarding any extra bits.

The Standard Pattern for Implementing Operations

When implementing new mathematical operations for BigFloat, a standard pattern emerges. Following this pattern helps ensure consistency and correctness across all operations.

First, you check for special cases, NaN, infinity, and zero, and handle them immediately. These cases often yield simple, well-defined results that require no computation. For instance, anything times NaN is NaN, and zero times anything is zero.

Next, you extract the raw significand and exponent values from your operands and prepare to work with them directly. You perform the mathematical operation on the significands using BigInt arithmetic, while tracking how it affects the exponents. If you're multiplying, the exponents add. If you're dividing, the exponents subtract. For addition and subtraction, you might need to align the numbers by adjusting exponents and shifting significands.

Finally, before returning your result, you call the `_normalize()` method to ensure that the result is in canonical form with the proper precision applied. This normalization step adjusts the exponent to bring the significand into the [1, 2) range, applies the requested rounding mode, and sets any necessary flags.

By following this consistent pattern, we preserve the invariants on which the rest of the codebase depends and make the code easier to understand and maintain.

How BigFloat Relates to IEEE 754

While BigFloat draws considerable inspiration from the IEEE 754 floating-point standard that underlies most modern computer arithmetic, it differs in several key ways that reflect its different design goals.

The most obvious difference is that IEEE 754 formats use fixed-width formats, whereas JavaScript variables have variable precision. BigFloat, by contrast, can have any precision you want. If you need 1000 decimal digits, you can have them.

IEEE 754 specifies precision in terms of bits, which makes sense for hardware implementation but is less intuitive for users. BigFloat uses decimal digits for precision specification, making it easier to request "I want 50 digits of accuracy" without converting to binary.

In IEEE 754 binary formats, there's an implicit leading 1 bit that isn't actually stored in the representation; the standard exploits the fact that normalized numbers always have a

leading 1 to save 1 bit of storage. BigFloat stores the entire significand explicitly, making the implementation more straightforward, even if it incurs a small space overhead.

Both systems support different rounding modes, but BigFloat can implement a broader range of rounding strategies with greater flexibility because it is not constrained by hardware.

Special values like NaN and infinity exist in both systems. Still, BigFloat handles them through explicit flags in the `_special` field rather than through special bit patterns in the exponent and significand fields as IEEE 754 does.

Despite these differences, the fundamental concepts are similar enough that anyone familiar with IEEE 754 floating-point arithmetic will find BigFloat's design intuitive. We get the efficiency and elegant mathematics of IEEE 754-style floating-point while adding the flexibility and unlimited precision of arbitrary-precision arithmetic libraries. This combination makes BigFloat both powerful and practical for a wide range of computational tasks requiring high precision.

Developing BigFloat with AI Assistance: A Collaborative Journey

The Initial Vision and Communication Challenge

When I set out to develop this new version of BigFloat, I had a clear vision of its internal structure. Having previously implemented an arbitrary-precision library in C++, I knew from experience that certain design decisions could make or break the project. I wanted to specify precision in decimal digits rather than bits, because that's how people naturally think about numerical accuracy, as previously described.

The internal representation I envisioned would use JavaScript's BigInt for the significand, paired with a binary exponent system similar to IEEE 754. This seemed like the perfect marriage of convenience and efficiency. BigInt would handle all the heavy lifting of arbitrary-precision integer arithmetic, while the binary exponent would keep scaling operations fast and simple.

When I began working with Claude AI to implement this vision, I discovered something fascinating and occasionally frustrating: a subtle yet critical gap between what I knew in my mind and what the AI initially understood from my descriptions. The fundamental challenge centered on how to interpret the BigInt significand with a hidden binary point after the first set bit.

The Conceptual Hurdle

On the surface, my requirements seemed straightforward enough. I explained that the significand should be a BigInt, the exponent should be binary, and the precision should be specified in decimal digits. I described how I wanted normalized numbers to always have their significand in the range [1, 2) when interpreted with the hidden binary point.

These weren't unusual requirements; they followed naturally from my experience with a previous C++ implementation and from understanding IEEE 754 floating-point formats.

However, the AI initially struggled to fully grasp what it meant to interpret a BigInt "with a hidden dot after the first set bit." This isn't just a minor detail; it's the conceptual foundation for how arithmetic operations need to work. When you multiply two numbers, you're not just multiplying their significands as plain integers and adding their exponents. You're multiplying two values that each represent something in the [1, 2) range, which produces a result in the [1, 4) range, and then you need to renormalize.

Early attempts to implement the basic arithmetic operations, addition, subtraction, and multiplication, revealed this misunderstanding. The AI would generate code that treated the significands as straightforward integers, performing the arithmetic correctly at that level but failing to account for where the hidden binary point should end up. The results would be off by a factor of 2, the normalization logic would be incorrect, or the exponent adjustments wouldn't adequately compensate for the significand manipulations.

The Breakthrough Moment

The turning point came when I invested time in really nailing down the internal format documentation and working through concrete examples step by step. Instead of just stating the rules abstractly, I walked through specific cases: "Here's how 2.5 is represented. The significand is this specific BigInt value, the exponent is this specific number, and here's exactly how you interpret them together to get 2.5."

I showed examples of what happens when you multiply 2.5 by 3.0. The significands multiply to yield a specific BigInt value. The exponents add. But now the product significand lies in a different range than [1, 2), so you need to shift it and adjust the exponent to bring it back into normalized form. Working through these concrete examples, with actual bit patterns and actual numbers, finally bridged the conceptual gap.

Once the AI truly grasped that the significand wasn't just an integer but an integer interpreted as a fixed-point number with the binary point in a specific position, everything clicked into place. Suddenly, implementing the basic operations became straightforward. The AI could see why the exponent needed to be adjusted in specific ways, why the significand needed to be shifted in certain circumstances, and how the normalization process should work.

Rapid Development After Understanding

With this foundational understanding established, the development pace accelerated dramatically. The basic arithmetic operations, addition, subtraction, and multiplication, were quickly developed and debugged. The AI could now reason correctly about what needed to happen at each step.

Division followed naturally once multiplication was working, because the inverse relationship was clear. Square root, which uses Newton's method applied to the reciprocal, converged smoothly because the AI understood how to track the binary point through iterative refinement. The transcendental functions, such as sine, cosine, and logarithm, were more complex in their mathematical algorithms. Still, the underlying representation issues were no longer a stumbling block, and I often just told the AI to generate the equivalent JavaScript function from the C++ code in my arbitrary-precision library.

What might have taken weeks of solo development happened in a much shorter timeframe through this collaborative process. The AI could quickly generate large blocks of code, allowing me to focus on reviewing the logic, testing edge cases, and refining the algorithms. When bugs appeared, and they certainly did, I could usually identify them quickly because we now shared a common understanding of how the representation worked. It certainly also helped that I could feed it large chunks of my C++ implementation, where the logic had been fully debugged, and ask it to convert them to equivalent JavaScript code.

The Iterative Refinement Process

Even after the initial breakthrough, development wasn't entirely smooth sailing. Several iterations were required to get the details right. For instance, handling negative numbers correctly required careful consideration of how signs propagate through operations. The value 0 required special handling because it doesn't fit the standard normalized representation. Infinity and NaN require separate logic paths.

The rounding modes presented another layer of complexity. It's one thing to implement "round to nearest" correctly, but quite another to handle all the different rounding modes (toward zero, away from zero, toward positive infinity, toward negative infinity) consistently across all operations. Each mode requires slightly different logic for deciding which way to round when you're truncating the significand to fit the target precision.

The AI proved invaluable during these refinement iterations. Once it understood the core representation, it could apply that understanding to address the various special cases. I could say, "Now we need to handle the case where both operands are infinity in multiplication," and the AI could reason through the mathematically correct answer and generate appropriate code.

Lessons About AI-Assisted Development

This experience taught me several valuable lessons about working with AI as a development partner. First and most importantly, the quality of the output depends critically on establishing a shared understanding of the fundamental concepts. Time invested in precise, detailed explanations of the core ideas pays enormous dividends later. Those early struggles to communicate the hidden binary point concept weren't wasted effort; they were essential groundwork.

Second, concrete examples are worth their weight in gold when working with AI. Abstract descriptions of algorithms or data structures can be misinterpreted in subtle ways. Walking through specific examples with actual numbers, showing precisely what bits go where and why, eliminates ambiguity and builds a reliable shared understanding.

Third, the AI excels at applying established patterns once it understands them. After we got the basic arithmetic operations working correctly, implementing additional mathematical functions became much faster because the AI could follow the same patterns. It understood how to handle special cases, maintain precision, normalize results, and structure the code consistently with what we'd already built.

Fourth, iterative development works exceptionally well in this collaborative model. Rather than trying to specify everything perfectly upfront, it's often better to build something, test it, identify what needs improvement, and then refine. The AI can quickly generate new versions incorporating feedback, making this iteration cycle much faster than traditional solo development.

The Value of Domain Knowledge

My prior experience implementing arbitrary precision arithmetic in C++ proved invaluable throughout this process. I knew what the finished product should look like. I understood the mathematical properties that needed to be preserved. I could recognize when something was wrong, even if I didn't immediately know how to fix it.

This domain knowledge allowed me to guide the AI effectively. When the generated code didn't work correctly, I could often identify the conceptual misunderstanding and explain it more clearly. When making design decisions, should we cache the bit length? How much extra precision should we use for intermediate calculations? I could draw on previous experience to make informed choices.

The collaboration worked because it combined the AI's ability to rapidly generate code with my ability to provide direction, validation, and domain expertise. Neither of us could have built BigFloat as quickly alone. The AI would have struggled with the conceptual foundation, and I would have spent far more time typing and debugging.

Reflections on the Process

Looking back at the development of BigFloat, I'm struck by how different this process was from traditional software development. The initial phase, where we worked to establish a shared understanding of the internal representation, felt almost like teaching. I was explaining concepts, providing examples, and checking for understanding, much as I might with a junior developer, except the "student" could generate hundreds of lines of code in seconds once the understanding was in place.

The middle phase, where we rapidly implemented the core functionality, felt more like pair programming with an incredibly fast typist who never got tired. I could focus on the

algorithmic and mathematical aspects while the AI handled the mechanical work of translating those ideas into code. When I spotted an issue or wanted to try a different approach, we could iterate quickly without the usual friction of rewriting large sections.

The later phases, which implemented more specialized functions and handled edge cases, demonstrated how the AI could extrapolate from established patterns. I didn't need to explain every detail of every function. Once we had a few examples working correctly, the AI could apply similar logic to new situations, with me providing guidance primarily when something truly novel arose.

What emerged from this collaborative process is a library that reflects both the careful design considerations I brought from my C++ experience and the rapid implementation capabilities enabled by AI assistance. The core architecture, the use of BigInt for the significand, the binary exponent system, and the decimal precision specification all came from my vision and experience. But the actual implementation, the hundreds of functions and thousands of lines of code, came together far more quickly than traditional development would have allowed.

This experience has convinced me that AI-assisted development, when done thoughtfully with clear communication and appropriate human guidance, can dramatically accelerate the creation of complex software systems. The key is recognizing that the AI is a tool that amplifies human expertise rather than replacing it. You need to understand what you're building, evaluate whether the generated code is correct, and guide the overall architecture and design. But within those constraints, AI assistance can transform what might have been a months-long project into something achievable in a much shorter timeframe.

BigFloat

Support for arbitrary-precision floating-point numbers in JavaScript

Constructor

```
new BigFloat(value, precision, mode)      // Invoked as a Constructor  
BigFloat(value, precision, mode)          // Invoked as a Conversion
```

Arguments

<i>value</i>	Optional Floating-point number, BigFloat object, or any JavaScript Number or String representing a number. If omitted, it is set to zero.
<i>precision</i>	Optional. The number of significant digits (Base 10) If the precision argument is omitted, it is treated as BigFloat.defaultPrecision, which has a default value of 20 digits. If <i>BigFloat</i> is invoked as a conversion, the value parameter is converted to a <i>BigFloat number</i> and returned.
<i>mode</i>	Optional rounding mode. If omitted, the default rounding mode is taken from BigFloat.defaultRounding, which has a default value of BigFloat.RoundingMode.NEAREST

Returns

Returns a BigFloat object initialized with the floating-point value. If *BigFloat* is invoked as a conversion, the *value* parameter is converted to a *BigFloat number* and returned. If *value* is another BigFloat number, then that is returned. If *the value is undefined*, it is treated as zero.

Also note that precision is an optional argument; if omitted, the default property BigFloat.defaultPrecision is used unless the value is a BigFloat object, in which case it inherits the precision from the BigFloat object.

Regardless of whether invoked as a new constructor or as a Conversion BigFloat constructor, it always returns a BigFloat normalized number.

Example:

```
x=New BigFloat(1.5);      // Return a new BigFloat object with precision BigFloat.precision (20);  
x=new BigFloat("1.5E2", 43); // Return a BigFloat object with precision of 43 digits  
y=new BigFloat(x);        // Return a new BigFloat object with an inherent precision of 43 digits  
y=new BigFloat(x,25);     // Return a new BigFloat object rounded down from 43 to 25 digits  
  
x=BigFloat(1.5);          // return a new BigFloat object with precision BigFloat.precision (20);  
x=BigFloat("1.5E2", 43); // Return a BigFloat object with precision of 43 digits  
y=BigFloat(x);           // Return a new BigFloat object with an inherent precision of 43 digits  
y=BigFloat(x,25);         // Return a new BigFloat object rounded down from 43 to 25 digits
```

BigFloat Internals.

A BigFloat object uses a binary floating-point representation inspired by IEEE 754, storing numbers in the form: $\text{value} = \text{sign} \times \text{significand} \times 2^{\text{exponent}}$. This structure enables efficient arithmetic operations while providing true arbitrary precision capabilities.

Object Members

Each BigFloat object contains the following internal members:

- **_sign:** The sign of the number: +1 for positive, -1 for negative
- **_exponent:** The binary exponent representing powers of 2 (the value scales the significand by 2^{exponent})
- **_significand:** The mantissa stored as a JavaScript BigInt, representing the significant digits in binary
- **_precision:** The number of decimal digits of precision for the fractional part of a normalized number
- **_rounding:** The rounding mode to use when normalizing results after operations
- **_special:** A flag indicating special values: null (normal number), "zero", "nan", "inf"
- **_bits:** The number of bits in the significand (cached for performance; recalculated if negative)

Normalized Representation

BigFloat stores numbers in normalized form, where the significand is in the range [1, 2). This is achieved by interpreting the significand with an implicit binary point after the first set bit. For example, if the significand is 0b1101 (13 in decimal), it is interpreted as 1.101_2 , which equals 1.625 in decimal.

The complete value is then: $\text{sign} \times 1.101_2 \times 2^{\text{exponent}}$. This normalization ensures that the leading bit of the significand is always 1 (except for 0), providing a unique representation for each number and maximizing the precision of the significand's bits.

Zero is a special case with no set bits in the significand and is handled via the `_special` flag.

BigInt for Performance

By using JavaScript's native BigInt type to store the significand, BigFloat achieves dramatic performance improvements. The heavy lifting for basic operations (+, -, ×) is performed by the JavaScript engine's highly optimized BigInt implementation (typically written in C++), rather than by interpreted JavaScript code that manipulates arrays of digits.

Binary Exponent System

The binary exponent (base 2) makes the structure similar to IEEE 754 floating-point. This choice provides several advantages: multiplying by powers of 2 is trivial (add the exponent), detecting powers of 2 is immediate, and many numerical algorithms naturally operate on binary scaling.

Implicit Binary Point

The significand is not interpreted as a plain integer but rather with an implicit binary point after the first bit. When performing multiplication, addition, or subtraction, the bit patterns produced are identical to those of integer operations; the only difference is tracking the location of the implicit binary point in the result. After each operation, normalization ensures the binary point is positioned correctly, maintaining the [1, 2) range for the significand.

Separate Sign Handling

Although JavaScript's BigInt supports two's complement representation for negative numbers, BigFloat uses a separate `_sign` field. This design choice simplifies the arithmetic logic and makes the implementation more intuitive, as operations can work with unsigned significands and handle signs separately.

Operation Guarantees

Every arithmetic and mathematical function in BigFloat guarantees that the result is returned as a normalized number. This means:

- The significand is normalized to the [1, 2) range (or marked as a special value)
- The exponent is adjusted to maintain the correct value
- The sign is set correctly
- Special flags (zero, NaN, infinity) are updated appropriately
- The result is rounded to the specified precision using the specified rounding mode
- The `_bits` are updated accordingly

Flexible Precision and Rounding

Unlike many arbitrary-precision libraries, BigFloat allows precision and rounding mode to be changed at runtime. Each BigFloat object carries its own precision and rounding mode settings, so you are not constrained to a fixed precision when creating variables. This flexibility is advantageous in complex calculations where you might want to increase precision for intermediate steps to minimize accumulated rounding errors, then reduce to the final target precision for the result. For example:

```
let x = new BigFloat(value, 50); // High precision for computation
let y = complexCalculation(x);
y._precision = 20;           // Reduce to target precision
y = y.normalize();          // Apply new precision
```

Constructor Design Philosophy

BigFloat deliberately avoids using JavaScript's class syntax for the constructor. Instead, it uses a traditional function constructor approach. This design allows for more flexible usage patterns, including implicit type conversion without the 'new' keyword:

```
let a = BigFloat(10.5);      // Works without "new"
let b = BigFloat.sqrt(a);    // Direct computation
let c = BigFloat(10.5).mul(2); // Chaining
```

This style resembles C++ cast operators or JavaScript's BigInt() constructor, providing a natural and convenient interface for users. The constructor automatically detects whether it was called with 'new' and handles both cases appropriately, ensuring consistent behavior regardless of how users instantiate BigFloat objects.

BigFloat's internal structure combines the performance benefits of native BigInt operations with the flexibility of per-instance precision control. The IEEE 754-inspired binary representation, implicit binary point, and automatic normalization work together to provide efficient, accurate, arbitrary-precision arithmetic. The design prioritizes both performance and usability, making BigFloat suitable for demanding numerical computations while remaining accessible and intuitive for developers.

Rounding modes

BigFloat supports handling multiple rounding modes per BigFloat object. The rounding mode comes into play when a new BigFloat number is assigned to a BigFloat object. Depending on the rounding mode, the variable is either rounded to the nearest (“NEAREST”) number of the BigFloat object precision, rounded down (“DOWN”), or rounded up. (“UP”) and finally rounding towards zero (“ZERO”). The rounding mode can be changed as needed using the Property method rounding().

Precision

Each BigFloat object has its own precision (number of fractional digits). The default precision is 20 decimal fractional digits. You can freely mix operations on BigFloat objects with different precisions. The operation always returns the maximum precision between the two numbers. E.g.

```
var x5=new BigFloat(1.23456,5); // create BigFloat with 5 decimal fraction digits
var xdefault=new BigFloat("1.2345678"); // create it with default precision (20digits)
var y= BigFloat.add(xdefault,x5); // y will be assigned with the precision of the
operation, which is the default precision in this case
x5.assign(y); // Assign x5 with y rounding to the precision of x5.
```

Since you can't overload the “=” operator in JavaScript, you have to use the method BigFloat.assign() to store the result with the precision of the left-hand side of the “=” operator. Otherwise, the left-hand side precision will be overwritten. E.g.

```
let a = new BigFloat(5.5,10);           // a has 10 digits precision
let b = new BigFloat(10.3,20);          // b has 20 digits precision
a.assign(BigFloat.add(a,b));            // the + is performed at 20 digits precision
// and store in a normalized form to 10 digits. Notice the left-hand side retains its
// precision and rounding mode.
a=BigFloat.add(a,b); // the + is performed at 20 digits precision and
```

```
// And a's precision and rounding mode is overwritten by the add operation result.
```

BigFloat.js Methods and Functions Reference

This document provides a comprehensive list of all methods and functions implemented in BigFloat.js. Functions are organized by category for easy reference.

Constructor

BigFloat(value, precision, rounding)

Creates a new BigFloat object from a number, string, BigInt, or another BigFloat. Precision defaults to BigFloat.defaultPrecision, rounding defaults to BigFloat.defaultRounding, and value to zero if omitted.

Configuration

BigFloat.defaultrounding(mode)

Get or set the default rounding mode

BigFloat.defaultprecision(prec)

Get or set the default precision

Instance Property Accessors

These methods get or set properties of a BigFloat instance:

- **sign(s)** - Get or set the sign (+1 or -1)
- **exponent(e)** - Get or set the binary exponent
- **precision(p)** - Get or set the precision in decimal digits
- **rounding(mode)** - Get or set the rounding mode

Type Checking Methods

- **isNaN()** - Returns true if the value is NaN
- **isFinite()** - Returns true if the value is finite (not NaN or infinity)
- **isInfinite()** - Returns true if the value is positive or negative infinity
- **isZero()** - Returns true if the value is zero
- **isInteger()** - Returns true if the value is an integer
- **isPositive()** - Returns true if the value is positive
- **isNegative()** - Returns true if the value is negative
- **isEven()** - Returns true if the value is an even integer
- **isOdd()** - Returns true if the value is an odd integer

Conversion and Formatting Methods

- **toNumber()** - Converts to a JavaScript Number (may lose precision)

- **toBigInt()** - Converts to a JavaScript BigInt Number (may lose precision)
- **toString(base)** - Converts to a string representation in the given base (default 10)
- **toExponential(digits)** - Returns exponential notation string
- **toFixed(fractionDigits)** - Returns fixed-point notation string
- **toPrecision(precision)** - Returns a string with specified precision

Basic Arithmetic Operations

All arithmetic operations are available as static methods. Some also have prototype methods.

- **BigFloat.add(a, b)** - Addition: $a + b$
- **BigFloat.sub(a, b)** - Subtraction: $a - b$
- **BigFloat.mul(a, b)** - Multiplication: $a \times b$
- **BigFloat.div(a, b)** - Division: $a \div b$
- **BigFloat.inverse(a)** - Reciprocal: $1/a$
- **BigFloat.sqrt(x)** - Square root: \sqrt{x}
- **BigFloat.abs(x)** - Absolute value: $|x|$

Comparison Operations

- **BigFloat.equal(a, b)** - Returns true if $a == b$
- **BigFloat.notequal(a, b)** - Returns true if $a \neq b$
- **BigFloat.less(a, b)** - Returns true if $a < b$
- **BigFloat.lessequal(a, b)** - Returns true if $a \leq b$
- **BigFloat.greater(a, b)** - Returns true if $a > b$
- **BigFloat.greaterequal(a, b)** - Returns true if $a \geq b$

Rounding and Integer Functions

- **BigFloat.floor(x)** - Rounds down to nearest integer (toward $-\infty$)
- **BigFloat.ceil(x)** - Rounds up to nearest integer (toward $+\infty$)
- **BigFloat.trunc(x)** - Truncates toward zero (removes fractional part)
- **BigFloat.round(x)** - Rounds to nearest integer (half away from zero)

Note: These functions also have prototype method versions: `x.floor()`, `x.ceil()`, `x.trunc()`, `x.round()`

Special Mathematical Functions

- **BigFloat.fmod(a, b)** - Floating-point remainder of a/b
- **BigFloat.modf(x)** - Returns [integer_part, fractional_part]
- **BigFloat.fma(a, b, c)** - Fused multiply-add: $a \times b + c$ with single rounding
- **BigFloat.frexp(x)** - Returns [mantissa, exponent] where $x = \text{mantissa} \times 2^{\text{exponent}}$
- **BigFloat.ldexp(x, exp)** - Returns $x \times 2^{\text{exp}}$

Adjacent Value Functions

- **succ()** - Returns the next representable number toward $+\infty$ (prototype method)
- **pred()** - Returns the next representable number toward $-\infty$ (prototype method)
- **BigFloat.nextafter(x, towards)** - Returns the next representable value of x in the direction of towards

Mathematical Constants

All constants are computed to arbitrary precision on demand:

- **BigFloat.PI(prec)** - Returns π to the specified precision
- **BigFloat.E(prec)** - Returns e (Euler's number) to the specified precision
- **BigFloat.LN2(prec)** - Returns $\ln(2)$ to the specified precision
- **BigFloat.LN5(prec)** - Returns $\ln(5)$ to the specified precision
- **BigFloat.LN10(prec)** - Returns $\ln(10)$ to the specified precision
- **BigFloat.SQRT2(prec)** - Returns $\sqrt{2}$ to the specified precision

Power and Exponential Functions

- **BigFloat.pow(x, y)** - Power: x^y (general exponent)
- **BigFloat.exp(x)** - Exponential: e^x
- **BigFloat.log(x) / BigFloat.ln(x)** - Natural logarithm: $\ln(x)$
- **BigFloat.log10(x)** - Base-10 logarithm: $\log_{10}(x)$

Note: pow(), exp() also have prototype method versions

Trigonometric Functions

- **BigFloat.sin(x)** - Sine of x (x in radians)
- **BigFloat.cos(x)** - Cosine of x (x in radians)
- **BigFloat.tan(x)** - Tangent of x (x in radians)
- **BigFloat.asin(x)** - Arcsine: $\sin^{-1}(x)$, returns value in $[-\pi/2, \pi/2]$
- **BigFloat.acos(x)** - Arccosine: $\cos^{-1}(x)$, returns value in $[0, \pi]$
- **BigFloat.atan(x)** - Arctangent: $\tan^{-1}(x)$, returns value in $(-\pi/2, \pi/2)$
- **BigFloat.atan2(y, x)** - Two-argument arctangent, returns angle in $[-\pi, \pi]$

Hyperbolic Functions

- **BigFloat.sinh(x)** - Hyperbolic sine
- **BigFloat.cosh(x)** - Hyperbolic cosine
- **BigFloat.tanh(x)** - Hyperbolic tangent
- **BigFloat.asinh(x)** - Inverse hyperbolic sine
- **BigFloat.acosh(x)** - Inverse hyperbolic cosine
- **BigFloat.atanh(x)** - Inverse hyperbolic tangent

Note: sinh(), cosh(), tanh() also have prototype method versions

Utility and Internal Methods

- **clone()** - Creates a deep copy of the BigFloat object
- **assign(a)** - Assigns the value of another BigFloat to this instance
- **dump2Console(label, verbose)** - Debug helper to print internal state to console

Rounding Modes

BigFloat supports four rounding modes:

- **BigFloat.RoundingMode.NEAREST** - Round to nearest (ties to even)
- **BigFloat.RoundingMode.ZERO** - Round toward zero (truncate)
- **BigFloat.RoundingMode.UP** - Round toward $+\infty$
- **BigFloat.RoundingMode.DOWN** - Round toward $-\infty$

Usage Examples

Basic arithmetic:

```
let a = new BigFloat("1.5", 50);
let b = new BigFloat("2.7", 50);
let sum = BigFloat.add(a, b);
let product = BigFloat.mul(a, b);
console.log(sum.toString());
```

Computing π to 100 digits:

```
let pi = BigFloat.PI(100);
console.log(pi.toString());
```

Trigonometric calculations:

```
let x = BigFloat.div(BigFloat.PI(50), new BigFloat(4)); //  $\pi/4$ 
let sinValue = BigFloat.sin(x);
console.log(sinValue.toString()); // Should be  $\approx 0.7071\dots$ 
```

Using fma for accurate computation:

```
let a = new BigFloat("1e20");
let b = new BigFloat("1e-20");
let c = new BigFloat("1");
// Computes  $a*b + c$  with single rounding
let result = BigFloat.fma(a, b, c);
```

Using `nextafter`:

```
let x = new BigFloat(1.5, 10);
let next = BigFloat.nextafter(x, new BigFloat(2));
console.log(next.toString()); // Next representable value > 1.5
```

BigFloat.js API Reference

Complete API Documentation

Version 3.0 - Arbitrary Precision Floating-Point Arithmetic Library

Overview

BigFloat.js provides arbitrary-precision floating-point arithmetic for JavaScript. This library allows you to perform mathematical operations with user-defined precision, going beyond the limitations of JavaScript's native Number type.

Constants and Configuration

Mathematical Constants

- `BigFloat.E(precision)` - Euler's number ($\approx 2.71828\dots$)
- `BigFloat.EPSILON(precision)` - Machine epsilon ($10^{-(precision)}$)

Returns the smallest positive value ϵ such that $1 + \epsilon \neq 1$ at the specified precision. Used for numerical comparisons, convergence testing, and tolerance-based equality checks.

Example: `const eps = BigFloat.EPSILON(50);
console.log(eps.toString()); // 1e-50`

- `BigFloat.PI(precision)` - Pi ($\approx 3.14159\dots$)
- `BigFloat.LN2(precision)` - Natural log of 2 ($\approx 0.69314\dots$)
- `BigFloat.LN10(precision)` - Natural log of 10 ($\approx 2.30258\dots$)

Instance Properties

- `x.EPSILON` - Machine epsilon for the instance's precision (read-only)

Convenience property that automatically returns `BigFloat.EPSILON(x._precision)` for the instance. Useful for tolerance-based comparisons where you want epsilon to match the instance's precision.

```
Example: const x = new BigFloat(1.0, 100); const eps = x.EPSILON;  
// Uses precision 100 // Equivalent to: const eps2 =  
BigFloat.EPSILON(100);
```

Special Value Constants

BigFloat.SPECIAL_NONE = 0	- Normal finite number
BigFloat.SPECIAL_ZERO = 1	- Zero (signed)
BigFloat.SPECIAL_INF = 2	- Infinity (signed)
BigFloat.SPECIAL_NAN = 3	- Not a Number

Rounding Mode Constants

BigFloat.ROUNDING_NEAREST = 0	- Round to nearest, ties to even
BigFloat.ROUNDING_UP = 1	- Round toward +infinity
BigFloat.ROUNDING_DOWN = 2	- Round toward -infinity
BigFloat.ROUNDING_ZERO = 3	- Round toward zero (truncate)

Default Settings

BigFloat.defaultPrecision - Get or set the default precision for new BigFloat objects

BigFloat.defaultRounding - Get or set the default rounding mode for new BigFloat objects

Methods and Functions

BigFloat.abs()

Return the absolute value of a BigFloat number

Synopsis: BigFloat.abs (x)

Arguments: x - The BigFloat number

Returns: The absolute value of x

Example:

```
var x = new BigFloat(-3.34); x=BigFloat.abs(x); // result  
3.34
```

BigFloat.acos()

Calculate the arc cosine of a BigFloat number

Synopsis: BigFloat.acos (x)

Arguments: x - The BigFloat number (must be in range [-1, 1])

Returns: The arc cosine of x in radians

Example:

```
var x = new BigFloat(0.5); BigFloat.acos(x); // result  
approx. 1.0471975511965976...
```

BigFloat.acosh()

Calculate the inverse hyperbolic cosine

Synopsis: BigFloat.acosh (x)

Arguments: x - The BigFloat number (must be ≥ 1)

Returns: The inverse hyperbolic cosine of x

Example:

```
var x = new BigFloat(2); BigFloat.acosh(x); // result approx.  
1.3169578969248166...
```

BigFloat.add()

Add two BigFloat numbers

Synopsis: BigFloat.add (a, b)

Arguments: a, b - The BigFloat numbers to be added

Returns: The sum of a and b. Precision is the maximum of the two operands.

Example:

```
var x = new BigFloat(1.2345); var y = new BigFloat(0.98765);  
BigFloat.add(x, y); // result 2.22215
```

BigFloat.asin()

Calculate the arc sine of a BigFloat number

Synopsis: BigFloat.asin(x)

Arguments: x - The BigFloat number (must be in range [-1, 1])

Returns: The arc sine of x in radians

Example:

```
var x = new BigFloat(0.75); BigFloat.asin(x); // result  
0.848062078981481...
```

BigFloat.asinh()

Calculate the inverse hyperbolic sine

Synopsis: BigFloat.asinh(x)

Arguments: x - The BigFloat number

Returns: The inverse hyperbolic sine of x

Example:

```
var x = new BigFloat(2); BigFloat.asinh(x); // result approx.  
1.4436354751788103...
```

BigFloat.assign()

Assign a value while preserving the precision and rounding mode of the target object

Synopsis: target.assign(source)

Arguments: source - The BigFloat to assign from

Returns: The target object with a new value but original precision

Example:

```
var x = new BigFloat(0.75, 10); var y = new BigFloat(0, 5);  
y.assign(x); // y = 0.75 with precision 5
```

BigFloat.atan()

Calculate the arc tangent of a BigFloat number

Synopsis: BigFloat.atan(x)

Arguments: x - The BigFloat number

Returns: The arc tangent of x in radians

Example:

```
var x = new BigFloat(0.75); BigFloat.atan(x); // result  
0.6435011087932844...
```

BigFloat.atan2()

Calculate the arc tangent of y/x with the correct quadrant

Synopsis: BigFloat.atan2(y, x)

Arguments: y - The Y coordinate x - The X coordinate

Returns: The angle in radians between $-\pi$ and $+\pi$

Example:

```
var x = new BigFloat(2), y = new BigFloat(1);
BigFloat.atan2(y, x); // result 0.463647609000806...
```

BigFloat.atanh()

Calculate the inverse hyperbolic tangent

Synopsis: BigFloat.atanh(x)

Arguments: x - The BigFloat number (must be in range (-1, 1))

Returns: The inverse hyperbolic tangent of x

Example:

```
var x = new BigFloat(0.5); BigFloat.atanh(x); // result
approx. 0.5493061443340548...
```

BigFloat.ceil()

Round up to the nearest integer

Synopsis: BigFloat.ceil(x) or x.ceil()

Arguments: x - The BigFloat number

Returns: The smallest integer greater than or equal to x

Example:

```
var x = new BigFloat(1.09); BigFloat.ceil(x); // result 2
```

BigFloat.clone()

Create a copy of a BigFloat object

Synopsis: x.clone()

Returns: A new BigFloat object with the same value, precision, and rounding mode

Example:

```
var x = new BigFloat(1.23, 10); var y = x.clone();
```

BigFloat.cos()

Calculate the cosine of a BigFloat number

Synopsis: BigFloat.cos (x)

Arguments: x - The BigFloat number (in radians)

Returns: The cosine of x

Example:

```
var x = new BigFloat(0.75); BigFloat.cos(x); // result  
0.731688868738209...
```

BigFloat.cosh()

Calculate the hyperbolic cosine

Synopsis: BigFloat.cosh (x) or x.cosh()

Arguments: x - The BigFloat number

Returns: The hyperbolic cosine of x

Example:

```
var x = new BigFloat(1); BigFloat.cosh(x); // result approx.  
1.5430806348152437...
```

BigFloat.defaultPrecision()

Get or set the default precision for new BigFloat objects

Synopsis: BigFloat.defaultPrecision (precision)

Arguments: precision - Optional. If provided, sets the new default precision

Returns: The current default precision

Example:

```
BigFloat.defaultPrecision(20); var x = new BigFloat(1.5); //  
x has precision 20
```

BigFloat.defaultRounding()

Get or set the default rounding mode for new BigFloat objects

Synopsis: BigFloat.defaultRounding (mode)

Arguments: mode - Optional. If provided, sets the new default rounding mode

Returns: The current default rounding mode

Example:

```
BigFloat.defaultRounding(BigFloat.ROUNDING_UP); var x = new  
BigFloat(1.5);
```

BigFloat.div()

Divide two BigFloat numbers

Synopsis: BigFloat.div(a, b)

Arguments: a - The dividend b - The divisor

Returns: The quotient a/b

Example:

```
var x = new BigFloat(3E+3); var y = new BigFloat(2);  
BigFloat.div(x, y); // result 1.5E+3
```

BigFloat.E()

Calculate Euler's number (e) to the specified precision

Synopsis: BigFloat.E(precision)

Arguments: precision - Optional. The decimal precision (default: defaultPrecision)

Returns: The value of e ($\approx 2.71828\dots$)

Example:

```
BigFloat.E(10); // result 2.7182818285
```

BigFloat.equal()

Test if two BigFloat numbers are equal

Synopsis: BigFloat.equal(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if a equals b, false otherwise

Example:

```
var x = new BigFloat(3.5); BigFloat.equal(x, x); // result  
true
```

BigFloat.EPSILON()

Calculate machine epsilon for the requested precision

Synopsis: BigFloat.EPSILON(precision)

Arguments: precision - The required decimal precision (number of decimal digits)

Returns: A BigFloat representing machine epsilon (ϵ) for the specified precision. Machine epsilon is the smallest positive value such that $1 + \epsilon \neq 1$ in floating-point arithmetic.

Description: For a given precision p (in decimal digits), machine epsilon is defined as:

$$\text{EPSILON} = 10^{-p}$$

This represents the smallest representable difference at the precision level. It is commonly used for:

- Testing numerical equality with tolerance
- Checking convergence of iterative algorithms
- Determining when computed results are "close enough"

Example:

```
// Get epsilon for 50 decimal digits const eps =
BigFloat.EPSILON(50); console.log(eps.toString()); // 1e-50
// Test equality with tolerance const a = new
BigFloat("1.0000000001", 50); const b = new BigFloat("1.0",
50); const diff = BigFloat.abs(BigFloat.sub(a, b)); if
(diff.lessThan(BigFloat.EPSILON(50))) {
  console.log("Values are equal within precision"); } else {
  console.log("Values differ"); } // Verify the epsilon
property:  $1 + \epsilon \neq 1$  const one = new BigFloat(1, 50); const
epsilon = BigFloat.EPSILON(50); const onePlusEps =
BigFloat.add(one, epsilon);
console.log(BigFloat.equal(onePlusEps, one)); // false
(correct!) // But  $1 + \epsilon/2$  should equal 1 (rounds back) const
halfEps = BigFloat.div(epsilon, new BigFloat(2, 50)); const
onePlusHalf = BigFloat.add(one, halfEps);
console.log(BigFloat.equal(onePlusHalf, one)); // true
```

Instance Property:

Machine epsilon is also available as an instance property that automatically uses the instance's precision:

```
const x = new BigFloat(1.0, 100); const eps = x.EPSILON; //
Automatically uses precision 100 // Equivalent to: const eps2
= BigFloat.EPSILON(100); console.log(BigFloat.equal(eps,
eps2)); // true
```

Comparison with Math.EPSILON:

JavaScript's Math.EPSILON is a fixed constant (approximately 2.22e-16) for IEEE 754 double precision. In contrast, BigFloat.EPSILON() computes epsilon for arbitrary precision:

```
// JavaScript double precision console.log(Math.EPSILON); //
2.220446049250313e-16 (fixed) // BigFloat arbitrary precision
console.log(BigFloat.EPSILON(10).toString()); // 1e-10
console.log(BigFloat.EPSILON(50).toString()); // 1e-50
console.log(BigFloat.EPSILON(1000).toString()); // 1e-1000
```

See Also

BigFloat.PI(), BigFloat.E(), x.EPSILON (instance property)

BigFloat.exp()

Calculate e raised to the power of x

Synopsis: BigFloat.exp(x) or x.exp()

Arguments: x - The BigFloat exponent

Returns: e^x

Example:

```
var x = new BigFloat(2.5); BigFloat.exp(x); // result approx  
12.182493960703473...
```

BigFloat.exponent()

Get or set the exponent of a BigFloat number

Synopsis: x.exponent(e)

Arguments: e - Optional. If provided, sets the exponent to this value

Returns: The exponent of the BigFloat number

Example:

```
var x = new BigFloat(1.2345E-6); x.exponent(); // result -6
```

BigFloat.floor()

Round down to the nearest integer

Synopsis: BigFloat.floor(x) or x.floor()

Arguments: x - The BigFloat number

Returns: The largest integer less than or equal to x

Example:

```
var x = new BigFloat(1.09); BigFloat.floor(x); // result 1
```

BigFloat.fma()

Fused multiply-add: compute $(a \times b) + c$ with no intermediate rounding

Synopsis: BigFloat.fma(a, b, c)

Arguments: a, b - The numbers to multiply, c - The number to add

Returns: The result of $(a \times b) + c$ with extended precision

Example:

```
var a = new BigFloat(2), b = new BigFloat(3), c = new  
BigFloat(4); BigFloat.fma(a, b, c); // result 10
```

BigFloat.fmod()

Calculate the floating-point remainder of a/b

Synopsis: BigFloat.fmod(a, b)

Arguments: a - The dividend b - The divisor

Returns: The floating-point remainder a - n×b, where n is the integer quotient

Example:

```
var a = new BigFloat(5.5), b = new BigFloat(2);
BigFloat.fmod(a, b); // result 1.5
```

BigFloat.frexp()

Extract mantissa and exponent (mantissa is in range [0.5, 1))

Synopsis: BigFloat.frexp(x)

Arguments: x - The BigFloat number

Returns: Object {mantissa, exponent} where x = mantissa × 2^{exponent}

Example:

```
var x = new BigFloat(12.5); var result = BigFloat.frexp(x); // 
mantissa ≈ 0.78125, exponent = 4
```

BigFloat.greater()

Test if a is greater than b

Synopsis: BigFloat.greater(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if a > b, false otherwise

Example:

```
var x = new BigFloat(3.5), y = new BigFloat(1.2);
BigFloat.greater(x, y); // result true
```

BigFloat.greaterequal()

Test if a is greater than or equal to b

Synopsis: BigFloat.greaterequal(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if a ≥ b, false otherwise

Example:

```
var x = new BigFloat(3.5); BigFloat.greaterequal(x, x); // 
result true
```

BigFloat.inverse()

Calculate the multiplicative inverse (1/x)

Synopsis: BigFloat.inverse(x)

Arguments: x - The BigFloat number (must not be zero)

Returns: 1/x

Example:

```
var x = new BigFloat(4); BigFloat.inverse(x); // result 0.25
```

[BigFloat.isEven\(\)](#)

Test if a number is even

Synopsis: x.isEven()

Returns: true if x is an even integer, false otherwise

Example:

```
var x = new BigFloat(4); x.isEven(); // result true
```

[BigFloat.isFinite\(\)](#)

Test if a number is finite

Synopsis: x.isFinite()

Returns: true if x is finite (not NaN or Infinity), false otherwise

Example:

```
var x = new BigFloat(1.5); x.isFinite(); // result true
```

[BigFloat.isInfinite\(\)](#)

Test if a number is infinite

Synopsis: x.isInfinite()

Returns: true if x is +Infinity or -Infinity, false otherwise

Example:

```
var x = new BigFloat(Infinity); x.isInfinite(); // result  
true
```

[BigFloat.isInteger\(\)](#)

Test if a number is an integer

Synopsis: x.isInteger()

Returns: true if x is an integer, false otherwise

Example:

```
var x = new BigFloat(5); x.isInteger(); // result true
```

BigFloat.isNaN()

Test if a number is Not-a-Number

Synopsis: `x.isnan()`

Returns: true if `x` is NaN, false otherwise

Example:

```
var x = new BigFloat(NaN); x.isnan(); // result true
```

BigFloat.isNegative()

Test if a number is negative

Synopsis: `x.isNegative()`

Returns: true if `x < 0`, false otherwise

Example:

```
var x = new BigFloat(-5); x.isNegative(); // result true
```

BigFloat.isOdd()

Test if a number is odd

Synopsis: `x.isOdd()`

Returns: true if `x` is an odd integer, false otherwise

Example:

```
var x = new BigFloat(5); x.isOdd(); // result true
```

BigFloat.isPositive()

Test if a number is positive

Synopsis: `x.isPositive()`

Returns: true if `x > 0`, false otherwise

Example:

```
var x = new BigFloat(5); x.isPositive(); // result true
```

BigFloat.isZero()

Test if a number is zero

Synopsis: `x.isZero()`

Returns: true if `x` is zero (positive or negative), false otherwise

Example:

```
var x = new BigFloat(0); x.isZero(); // result true
```

BigFloat.ldexp()

Multiply a number by a power of 2

Synopsis: BigFloat.ldexp(x, exp)

Arguments: x - The BigFloat number exp - The exponent (integer)

Returns: $x \times 2^{\text{exp}}$

Example:

```
var x = new BigFloat(1.5); BigFloat.ldexp(x, 3); // result 12
```

BigFloat.less()

Test if a is less than b

Synopsis: BigFloat.less(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if $a < b$, false otherwise

Example:

```
var x = new BigFloat(3.5), y = new BigFloat(1.2);
BigFloat.less(y, x); // result true
```

BigFloat.lessequal()

Test if a is less than or equal to b

Synopsis: BigFloat.lessequal(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if $a \leq b$, false otherwise

Example:

```
var x = new BigFloat(3.5); BigFloat.lessequal(x, x); // result true
```

BigFloat.LN2()

Calculate the natural logarithm of 2

Synopsis: BigFloat.LN2(precision)

Arguments: precision - Optional. The decimal precision (default: defaultPrecision)

Returns: $\ln(2) \approx 0.693147\dots$

Example:

```
BigFloat.LN2(10); // result 0.6931471805...
```

BigFloat.LN10()

Calculate the natural logarithm of 10

Synopsis: BigFloat.LN10 (precision)

Arguments: precision - Optional. The decimal precision (default: defaultPrecision)

Returns: $\ln(10) \approx 2.302585\dots$

Example:

```
BigFloat.LN10(10); // result 2.3025850929...
```

BigFloat.log()

Calculate the natural logarithm (alias: ln)

Synopsis: BigFloat.log(x) or BigFloat.ln(x)

Arguments: x - The BigFloat number (must be > 0)

Returns: The natural logarithm of x

Example:

```
var x = new BigFloat(2.5); BigFloat.log(x); // result approx.  
0.916290732...
```

BigFloat.log10()

Calculate the base-10 logarithm

Synopsis: BigFloat.log10(x)

Arguments: x - The BigFloat number (must be > 0)

Returns: The base-10 logarithm of x

Example:

```
var x = new BigFloat(100); BigFloat.log10(x); // result 2
```

BigFloat.modf()

Extract integer and fractional parts

Synopsis: BigFloat.modf(x) or x.modf()

Description: Extract the integer and fractional parts of a BigFloat number. The function splits a number into two components:

- **Integer part:** The whole number portion (truncated towards zero)
- **Fractional part:** The remainder after removing the integer part

Both parts retain the same sign as the original number, ensuring that:

$$x = \text{integerPart} + \text{fractionalPart}$$

Arguments: x - The BigFloat number

Returns: An array [integerPart, fractionalPart] where:

- integerPart - A BigFloat containing the integer portion (equivalent to trunc(x))

- `fractionalPart` - A BigFloat containing the fractional portion (equivalent to $x - \text{trunc}(x)$)
- Both parts have the same sign as x

Note: Returns an array with two BigFloat elements, not an object.`modf()` does not follow the C++ library standard since it returns an array with two elements.

Example:

Basic Usage

```
var x = new BigFloat(12.75); var result = BigFloat.modf(x); //  
result[0] = 12 (integer part) // result[1] = 0.75 (fractional part)  
// Using array destructuring (modern JavaScript): var [intPart,  
fracPart] = BigFloat.modf(x); console.log(intPart); // 12  
console.log(fracPart); // 0.75
```

Negative Numbers (both parts keep the sign)

```
var x = new BigFloat(-12.75); var [intPart, fracPart] =  
BigFloat.modf(x); console.log(intPart); // -12 (negative)  
console.log(fracPart); // -0.75 (negative) // Verify: x = intPart  
+ fracPart // -12.75 = -12 + (-0.75) ✓
```

Edge Cases

```
// Pure integer var x = new BigFloat(5); var [intPart, fracPart] =  
BigFloat.modf(x); // intPart = 5, fracPart = 0 // Pure fraction var  
x = new BigFloat(0.25); var [intPart, fracPart] = BigFloat.modf(x);  
// intPart = 0, fracPart = 0.25 // Zero var x = new BigFloat(0);  
var [intPart, fracPart] = BigFloat.modf(x); // intPart = 0, fracPart  
= 0
```

Difference from C++ `std::modf()`

BigFloat.js (JavaScript):

```
// Returns an ARRAY var [intPart, fracPart] = BigFloat.modf(x);
```

C++ `std::modf()`:

```
// Uses OUTPUT PARAMETER (pointer) double fracPart = std::modf(x, &intPart);  
// Returns frac, stores int in pointer
```

Key Differences:

Feature	BigFloat.js	C++ <code>std::modf</code>
Return type	Array [int, frac]	Fractional part only
Integer part	First array element	Via output parameter (pointer)
Fractional part	Second array element	Return value
Sign behavior	Both parts have same sign as x	Both parts have same sign as x
Usage pattern	<code>var [i, f] = modf(x)</code>	<code>double f = modf(x, &i)</code>

Why the difference?

- JavaScript doesn't have output parameters (no pointers)
- Returning an array is the idiomatic JavaScript approach

- Modern JavaScript array destructuring makes this elegant

BigFloat.mul()

Multiply two BigFloat numbers

Synopsis: BigFloat.mul(a, b)

Arguments: a, b - The BigFloat numbers to multiply

Returns: The product a × b

Example:

```
var x = new BigFloat(3), y = new BigFloat(2.5);
BigFloat.mul(x, y); // result 7.5
```

BigFloat.neg()

Return the absolute value of a BigFloat number

Synopsis: BigFloat.neg(x)

Arguments: x - The BigFloat number

Returns: The negated value of x (-x)

Example:

```
var x = new BigFloat(-3.34); x=BigFloat.neg(x); // result
3.34
```

BigFloat.nextafter()

Get the next representable number in a given direction

Synopsis: BigFloat.nextafter(x, towards)

Arguments: x - The starting BigFloat number, towards - The direction to move

Returns: The next representable value after x in the direction of towards

Example:

```
var x = new BigFloat(1.5, 10); var y = new BigFloat(2);
BigFloat.nextafter(x, y);
```

BigFloat.notequal()

Test if two numbers are not equal

Synopsis: BigFloat.notequal(a, b)

Arguments: a, b - The BigFloat numbers to compare

Returns: true if a ≠ b, false otherwise

Example:

```
var x = new BigFloat(3.5), y = new BigFloat(1.2);
BigFloat.notequal(x, y); // result true
```

BigFloat.PI()

Calculate π to the specified precision

Synopsis: BigFloat.PI(precision)

Arguments: precision - Optional. The decimal precision (default: defaultPrecision)

Returns: The value of π ($\approx 3.14159\dots$)

Example:

```
BigFloat.PI(10); // result 3.1415926536
```

BigFloat.pow()

Raise a number to a power

Synopsis: BigFloat.pow(x, y) or x.pow(y)

Arguments: x - The base y - The exponent

Returns: x^y

Example:

```
var x = new BigFloat(2.5), y = new BigFloat(3);
BigFloat.pow(x, y); // result 15.625
```

BigFloat.precision()

Get or set the precision of a BigFloat number

Synopsis: x.precision(p)

Arguments: p - Optional. If provided, sets the precision to this value

Returns: The precision (number of decimal digits)

Example:

```
var x = new BigFloat(1.5, 10); x.precision(); // result 10
x.precision(20); // changes precision to 20
```

BigFloat.pred()

Get the previous representable number (towards $-\infty$)

Synopsis: x.pred()

Returns: The next representable value towards negative infinity

Example:

```
var x = new BigFloat(1.5, 10); x.pred(); // slightly less
than 1.5
```

BigFloat.round()

Round to the nearest integer

Synopsis: BigFloat.round(x) or x.round()

Arguments: x - The BigFloat number

Returns: The nearest integer (0.5 rounds up)

Example:

```
var x = new BigFloat(1.5); BigFloat.round(x); // result 2
```

BigFloat.rounding()

Get or set the rounding mode

Synopsis: x.rounding(mode)

Arguments: mode - Optional. If provided, sets the rounding mode

Returns: The current rounding mode

Example:

```
var x = new BigFloat(1.5); x.rounding(); // result  
ROUNDING_NEAREST x.rounding(BigFloat.ROUNDING_UP);
```

BigFloat.sign()

Get or set the sign of a number

Synopsis: x.sign(s)

Arguments: s - Optional. If provided, sets the sign to this value

Returns: 1 for positive, -1 for negative, 0 for zero

Example:

```
var x = new BigFloat(-5); x.sign(); // result -1
```

BigFloat.sin()

Calculate the sine of a number

Synopsis: BigFloat.sin(x)

Arguments: x - The BigFloat number (in radians)

Returns: The sine of x

Example:

```
var x = new BigFloat(0.75); BigFloat.sin(x); // result  
0.681638760023334...
```

BigFloat.sinh()

Calculate the hyperbolic sine

Synopsis: BigFloat.sinh(x) or x.sinh()

Arguments: x - The BigFloat number

Returns: The hyperbolic sine of x

Example:

```
var x = new BigFloat(1); BigFloat.sinh(x); // result approx.  
1.1752011936438014...
```

[BigFloat.sqrt\(\)](#)

Calculate the square root

Synopsis: BigFloat.sqrt(x) or x.sqrt()

Arguments: x - The BigFloat number (must be ≥ 0)

Returns: The square root of x

Example:

```
var x = new BigFloat(2); BigFloat.sqrt(x); // result  
1.41421356237...
```

[BigFloat.SQRT2\(\)](#)

Calculate $\sqrt{2}$ to the specified precision

Synopsis: BigFloat.SQRT2(precision)

Arguments: precision - Optional. The decimal precision (default: defaultPrecision)

Returns: The value of $\sqrt{2}$ ($\approx 1.414213\dots$)

Example:

```
BigFloat.SQRT2(10); // result 1.4142135623
```

[BigFloat.sub\(\)](#)

Subtract two BigFloat numbers

Synopsis: BigFloat.sub(a, b)

Arguments: a, b - The BigFloat numbers

Returns: The difference a - b

Example:

```
var x = new BigFloat(5), y = new BigFloat(3); BigFloat.sub(x,  
y); // result 2
```

[BigFloat.succ\(\)](#)

Get the next representable number (towards $+\infty$)

Synopsis: x.succ()

Returns: The next representable value towards positive infinity

Example:

```
var x = new BigFloat(1.5, 10); x.succ(); // slightly greater  
than 1.5
```

BigFloat.tan()

Calculate the tangent of a number

Synopsis: BigFloat.tan(x)

Arguments: x - The BigFloat number (in radians)

Returns: The tangent of x

Example:

```
var x = new BigFloat(0.75); BigFloat.tan(x); // result  
0.931596459944348...
```

BigFloat.tanh()

Calculate the hyperbolic tangent

Synopsis: BigFloat.tanh(x) or x.tanh()

Arguments: x - The BigFloat number

Returns: The hyperbolic tangent of x

Example:

```
var x = new BigFloat(1); BigFloat.tanh(x); // result approx.  
0.7615941559557649...
```

BigFloat.toInt()

Convert to a BigInt (truncates towards zero)

Synopsis: x.toInt()

Returns: A BigInt with the integer value of x

Example:

```
var x = new BigFloat(12345.6789); x.toInt(); // result  
12345n
```

BigFloat.toExponential()

Format as exponential notation

Synopsis: x.toExponential(digits)

Arguments: digits - Optional. Number of digits after the decimal point

Returns: String in exponential notation (e.g., '1.23e+4')

Example:

```
var x = new BigFloat(12345.6789); x.toExponential(2); //  
result '1.23e+4'
```

BigFloat.toFixed()

Format as fixed-point notation

Synopsis: `x.toFixed(digits)`

Arguments: digits - Optional. Number of digits after decimal point (default: 0)

Returns: String in fixed-point notation

Example:

```
var x = new BigFloat(12345.6789); x.toFixed(2); // result  
'12345.68'
```

BigFloat.toNumber()

Convert to a JavaScript Number

Synopsis: `x.toNumber()`

Returns: A JavaScript Number (may lose precision)

Example:

```
var x = new BigFloat(1.23, 20); x.toNumber(); // result 1.23
```

BigFloat.toPrecision()

Format with specified significant digits

Synopsis: `x.toPrecision(digits)`

Arguments: digits - Optional. Number of significant digits

Returns: String with specified precision

Example:

```
var x = new BigFloat(12345.6789); x.toPrecision(3); // result  
'1.23e+4'
```

BigFloat.toString()

Convert to string representation

Synopsis: `x.toString(radix)`

Arguments: radix - Optional. Base (2, 10, or 16). Default is 10.

Returns: String representation in scientific notation

Example:

```
var x = new BigFloat(12345.6789); x.toString(); // result  
'1.23456789e+4'
```

BigFloat.trunc()

Truncate to integer (remove fractional part)

Synopsis: BigFloat.trunc(x) or x.trunc()

Arguments: x - The BigFloat number

Returns: The integer part (rounded towards zero)

Example:

```
var x = new BigFloat(1.9); BigFloat.trunc(x); // result 1
```

[parseBigFloat\(\)](#)

Parse a string to create a BigFloat

Synopsis: parseBigFloat(s, precision, rounding)

Arguments: s - The string to parse. precision - Optional. The precision rounding - Optional. The rounding mode

Returns: A new BigFloat object, or NaN if parsing fails

Example:

```
var x = parseBigFloat('1.23e-5'); var y =  
parseBigFloat('3.14', 20);
```

[BigFloat.toFixed\(\)](#)

Format using fixed-point notation

Synopsis: BigFloat_object.toFixed(digits)

Arguments: digits - Optional. Number of digits after the decimal point

Returns: String in fixed-point notation.

Example:

```
var x = new BigFloat(12345.6789); x.toFixed(2); // result  
'12345.68'
```

[BigFloat.toInteger\(\)](#)

Convert to an integer BigFloat

Synopsis: BigFloat_object.toInteger()

Returns: A BigFloat with integer value (truncated towards zero).

Example:

```
var x = new BigFloat(12345.6789); x.toInteger();
```

[BigFloat.toNumber\(\)](#)

Convert to a JavaScript Number

Synopsis: BigFloat_object.toNumber()

Returns: A JavaScript Number (may lose precision).

Example:

```
var x = new BigFloat(1.23456789, 20); x.toNumber(); //  
result 1.23456789
```

[BigFloat.toPrecision\(\)](#)

Format with significant digits

Synopsis: BigFloat_object.toPrecision(digits)

Arguments: digits - Optional. Number of significant digits

Returns: String with specified significant digits.

Example:

```
var x = new BigFloat(12345.6789); x.toPrecision(5); //  
result '12346'
```

[BigFloat.toString\(\)](#)

Convert to string representation

Synopsis: BigFloat_object.toString(radix)

Arguments: radix - Optional. Base (2, 10, or 16). Default is 10.

Returns: String representation in scientific notation.

Example:

```
var x = new BigFloat(12345.6789); x.toString(); // result  
'1.23456789e+4'
```

[BigFloat.trunc\(\)](#)

Truncate to an integer by removing the fractional part

Synopsis: BigFloat.trunc(x)

Arguments: x - The BigFloat number

Returns: Integer part (rounded towards zero).

Example:

```
var x = new BigFloat(1.09); BigFloat.trunc(x); // result 1
```

Appendix

Comparison of BigFloat with BigNumber.js and Decimal.js

This document provides a comprehensive comparison of BigFloat against two popular JavaScript arbitrary-precision libraries: BigNumber.js and Decimal.js. Both competing libraries are mature, well-tested projects by Michael McLaughlin that have been widely adopted in the JavaScript ecosystem. Understanding how BigFloat differs from these established libraries helps illuminate the unique design decisions and trade-offs inherent in each approach.

Part 1: BigFloat vs BigNumber.js

Internal Representation

The most fundamental difference between BigFloat and BigNumber.js lies in their internal representations, how they actually store and manipulate numbers at the bit level.

BigNumber.js Structure

BigNumber.js uses a decimal-based representation with three core fields:

- **s:** The sign (1 for positive, -1 for negative)
- **e:** The decimal exponent (base 10)
- **c:** The coefficient; an array of integers, each storing up to 14 decimal digits

The coefficient array breaks the number into chunks. For example, the number 1234567890123456789 would be stored in the coefficient array as [12345, 67890123456789], with each element representing a group of decimal digits. The constant `BASE = 1e14` defines this chunking; each array element can hold up to 14 decimal digits. This design allows BigNumber.js to perform arithmetic using JavaScript's native number type for intermediate calculations within each chunk, while the array structure provides unlimited precision overall.

The decimal exponent `e` represents powers of 10. So internally, BigNumber.js represents a value as:

$$\text{sign} \times \text{coefficient} \times 10^{\text{exponent}}$$

This decimal-centric approach makes sense from a human perspective; we think in decimals, we input decimal strings, and we expect decimal output. By staying in base 10 throughout, BigNumber.js avoids any conversion overhead between binary and decimal representations.

BigFloat Structure

BigFloat takes a radically different approach with a binary representation:

- **_sign:** The sign (1 for positive, -1 for negative)
- **_exponent:** The binary exponent (base 2, representing powers of 2)
- **_significand:** A single BigInt value storing the mantissa in binary
- **_precision:** Decimal precision specification
- **_bits:** Cached bit length of the significand
- **_rounding:** Rounding mode
- **_special:** Flag for special values (zero, NaN, infinity)

The significand is interpreted with an implicit binary point after the first set bit, representing a value in the normalized range [1, 2). This mirrors the IEEE 754 floating-point format used by hardware. The binary exponent scales this significand by powers of 2.

$$\text{value} = \text{sign} \times \text{significand} \times 2^{\text{exponent}}$$

By using a single BigInt for the significand rather than an array of decimal chunks, BigFloat gains several advantages. Operations like multiplication and bit-shifting become simpler because they operate on a single cohesive value. The binary exponent system makes scaling operations trivial; multiplying by powers of 2 is just exponent addition. Detecting powers of 2 is immediate. The normalized representation ensures every number has a canonical form.

Performance Implications

These different internal representations lead to distinct performance characteristics that manifest across various operations.

Arithmetic Operations

In BigNumber.js, arithmetic operations work with the decimal coefficient array. Addition and subtraction align the exponents, then process the coefficient arrays element by element, handling carries between chunks. Multiplication uses grade-school multiplication across array elements, with careful carry management. While this decimal approach is intuitive, it requires array manipulation and carries, which incur overhead. BigFloat's approach is different. Because the significand is a single BigInt, operations like multiplication are performed as single BigInt multiplications; the JavaScript engine handles all the complexity internally with highly optimized algorithms (Karatsuba, Toom-Cook, or even FFT-based multiplication for very large numbers). The binary exponent adds or subtracts. There's no array manipulation, no manual carry handling, and no chunking logic to manage.

For pure multiplication and division, BigFloat's approach tends to be faster because it leverages native BigInt performance. The engine's multiplication algorithms are implemented in highly optimized C++ or even assembly code, using techniques that pure JavaScript cannot access.

Conversion Overhead

However, BigFloat pays a price when converting between its internal binary representation and the decimal strings that users work with. When you create a BigFloat from a decimal string like "3.14159", the library must convert this base-10 input into a binary significand and exponent. Similarly, when converting back to a decimal string for display, there's an overhead cost for the binary-to-decimal conversion.

BigNumber.js avoids this conversion cost entirely because it works in decimal throughout. When you construct a BigNumber from "3.14159", it parses the digits directly into its decimal coefficient array with minimal processing. Outputting the number as a string is equally straightforward since the internal representation is already decimal.

The performance trade-off depends on usage patterns. If you're performing many arithmetic operations on each number (calculating complex mathematical functions, for instance), BigFloat's faster arithmetic operations outweigh the conversion overhead. If you're doing relatively few operations per number with frequent conversion between strings and numbers, BigNumber.js's decimal approach may be more efficient.

Precision Handling

Both libraries specify precision in decimal digits, which aligns with how humans think about numerical accuracy. However, they handle precision quite differently in practice.

BigNumber.js Approach

BigNumber.js uses a global configuration model. You set DECIMAL_PLACES as a library-wide setting that affects operations involving division and other functions that might produce infinite decimal expansions. The precision is specified in decimal places, and the internal decimal representation naturally aligns with this specification. When an operation needs to round, it simply truncates or rounds the decimal coefficient array at the appropriate position.

This global precision model is simple and works well for many applications. You configure the library once with your desired precision, and all subsequent operations respect that setting. The downside is less flexibility; if different parts of your calculation need different precision levels, you must reconfigure the library between operations.

BigFloat Approach

BigFloat takes a different approach by storing precision as an instance variable on each number. Every BigFloat object carries its own `_precision` field, allowing different numbers to have different precision levels simultaneously. You can increase precision for sensitive intermediate calculations and reduce it for final results, all within the same computation, without global reconfiguration.

This per-instance precision offers greater flexibility but adds complexity. The library must track precision through operations; when you multiply two BigFloats with different precisions, the result's precision must be determined intelligently. BigFloat typically uses the maximum precision of the operands or a configurable default, giving users fine-grained control.

Because BigFloat's internal representation is binary, there's a subtle conversion in precision semantics. The `_precision` field specifies decimal digits, but this translates to approximately $\text{_precision} \times 3.32$ bits internally (since $\log_2(10) \approx 3.32$). The library must maintain enough binary digits to guarantee the requested decimal precision, which involves careful rounding and truncation logic.

Mathematical Function Coverage

The range of mathematical functions provided by each library differs significantly.

BigNumber.js Functions

BigNumber.js provides the core arithmetic operations (addition, subtraction, multiplication, division) along with a more limited set of mathematical functions. It includes:

- Square root (`sqrt`)
- Power function (`pow`)
- Absolute value, negation
- Rounding and truncation functions
- Comparison operations
- Conversion to various formats

Notably absent are the transcendental functions: sine, cosine, logarithm, exponential, and their inverse and hyperbolic variants. For applications needing only financial calculations or basic arithmetic with arbitrary precision, this limitation isn't a problem. The smaller function set keeps the library compact and focused.

BigFloat Functions

BigFloat aims for comprehensive mathematical coverage, including all the functions that BigNumber.js provides, plus an extensive set of transcendental functions:

- Trigonometric: `sin`, `cos`, `tan`
- Inverse trigonometric: `asin`, `acos`, `atan`, `atan2`
- Hyperbolic: `sinh`, `cosh`, `tanh`
- Inverse hyperbolic: `asinh`, `acosh`, `atanh`
- Exponential and logarithmic: `exp`, `log`, `log10`, `ln`
- Power and root functions: `pow`, `sqrt`
- Special functions: `PI`, `LN(2)`, `E`, constant calculation

This comprehensive feature set makes BigFloat suitable for scientific computing, numerical analysis, and any application that requires advanced mathematical operations with arbitrary precision. The trade-off is a larger codebase and greater complexity.

API Design and Usability

The way users interact with these libraries reflects different design philosophies.

BigNumber.js API

BigNumber.js uses method chaining extensively, which allows for very readable code:

```
let result = new BigNumber('123.456')
  .plus('78.9')
  .times(2)
  .dividedBy(3)
  .toFixed(4);
```

The library provides both long method names (`absoluteValue`, `multipliedBy`) and short aliases (`abs`, `times`) to suit different coding styles. Methods return new BigNumber instances, making the library naturally immutable; operations don't modify the original number.

Configuration is set through the `BigNumber.config()` method, which sets global defaults that affect all BigNumber instances. This centralized configuration is simple and works well when you want consistent behavior across your application.

BigFloat API

BigFloat uses static methods for operations, following a pattern more similar to Math object conventions:

```
let a = new BigFloat('123.456');
let b = new BigFloat('78.9');
let result = BigFloat.div(
  BigFloat.mul(BigFloat.add(a, b), new BigFloat(2)),
  new BigFloat(3)
);
```

This functional style makes the operation explicit; you can see at a glance that `div`, `mul`, and `add` are BigFloat operations rather than object methods. The downside is less fluent chaining and more verbose code. However, this verbosity can improve clarity in complex expressions.

BigFloat's per-instance precision and rounding settings provide more flexibility but require more management. You can adjust these properties on individual numbers as needed rather than globally reconfiguring the library.

Memory and Performance Characteristics

Memory Footprint

BigNumber.js stores numbers as objects with a sign, exponent, and coefficient array. The array grows with the number of significant digits, with each element holding up to 14 decimal digits. For a 100-digit number, the coefficient array would have about eight elements ($100 \div 14 \approx 8$), each storing a JavaScript number.

BigFloat stores the significand as a single BigInt, which internally uses a variable-length array of 64-bit limbs (on 64-bit systems). A 100-digit decimal number requires

approximately 332 bits ($100 \times \log_2(10) \approx 332$), which fits in about 6 BigInt limbs. The memory overhead is comparable, though BigFloat includes additional fields, such as a bits cache and a `_special` flag.

In practice, memory usage is similar for both libraries, with differences usually negligible compared to the actual digit storage requirements.

Speed Comparisons

Performance varies significantly depending on the operation:

- **Construction from strings:** BigNumber.js is typically faster due to direct decimal parsing
- **Arithmetic (multiply, divide):** BigFloat tends to be faster, especially for large numbers, due to optimized BigInt operations
- **Square root:** BigFloat's division-free Newton's method is faster
- **Transcendental functions:** BigFloat provides these; BigNumber.js doesn't, so no comparison is possible
- **String conversion:** BigNumber.js is faster since no base conversion is needed

The performance picture is nuanced. For applications that frequently convert between strings and numbers with relatively few operations per number, BigNumber.js's decimal approach excels. For numerical computing with many operations per number, BigFloat's binary arithmetic advantage dominates.

Error Handling and Edge Cases

Both libraries handle special values (NaN, Infinity, and zero) and edge cases, but with different approaches.

BigNumber.js Approach

BigNumber.js uses null values in the coefficient array (`c`) to represent Infinity, with the sign indicating whether it is positive or negative infinity. NaN is defined similarly. This works, but it feels somewhat ad hoc; you're repurposing a field meant for digits to encode special states.

The library has configurable limits (`MIN_EXP` and `MAX_EXP`) that determine when underflow to zero or overflow to infinity occurs. These limits can be adjusted based on application needs.

BigFloat Approach

BigFloat uses an explicit `_special` field that can be null (for normal numbers), 'zero', 'inf', or 'nan'. This makes special value handling more explicit and easier to reason about. The code can check the `_special` field at the start of operations and handle special cases before attempting arithmetic.

Both approaches work fine in practice, but BigFloat's explicit special-value flag makes the code's intent clearer and special-case handling more systematic.

Part 2: BigFloat vs Decimal.js

Decimal.js, also by Michael McLaughlin, is actually a more feature-rich successor to BigNumber.js. It shares the same author and many of the same design principles while offering significantly expanded functionality. Comparing BigFloat to Decimal.js reveals both similarities and interesting differences in how comprehensive arbitrary-precision libraries can be designed.

Internal Representation

Decimal.js Structure

Like BigNumber.js, Decimal.js uses a decimal-based representation, though with some refinements:

- **s:** The sign (1 for positive, -1 for negative)
- **e:** The decimal exponent (base 10)
- **d:** The digits; an array of integers, where each element can hold up to 7 decimal digits

Notice that Decimal.js uses $\text{BASE} = 1e7$ (7 digits per element) rather than BigNumber.js's $1e14$ (14 digits per element). This choice trades some memory efficiency for potentially better performance in certain operations. With smaller chunks, intermediate calculations are less likely to overflow JavaScript's 53-bit integer precision limits, allowing more operations to use fast integer arithmetic without falling back to slower paths.

The fundamental representation remains decimal-based, sharing the same advantages (no conversion overhead for decimal I/O) and disadvantages (potentially slower arithmetic for large numbers) as BigNumber.js.

Mathematical Function Coverage

This is where Decimal.js truly distinguishes itself from its predecessor BigNumber.js. Decimal.js provides comprehensive mathematical function coverage that rivals BigFloat.

Decimal.js Functions

Decimal.js includes an extensive set of mathematical functions:

- Basic arithmetic: addition, subtraction, multiplication, division
- Trigonometric functions: sin, cos, tan
- Inverse trigonometric: asin, acos, atan, atan2
- Hyperbolic functions: sinh, cosh, tanh
- Inverse hyperbolic: asinh, acosh, atanh
- Exponential and logarithmic: exp, ln, log (arbitrary base), log2, log10
- Power and root functions: pow, sqrt, cbrt (cube root)
- Special functions and utilities: hypot, random, various rounding modes

This comprehensive coverage makes Decimal.js suitable for scientific computing and numerical analysis, the same domain as BigFloat. Both libraries recognize that arbitrary precision is most valuable when you can actually use it for advanced mathematical operations, not just basic arithmetic.

Implementation Approaches

Where BigFloat and Decimal.js differ most is in how they implement these mathematical functions.

Decimal.js implements transcendental functions using decimal arithmetic throughout. For example, its sine function uses a Taylor series expansion, but all the arithmetic is performed in the decimal coefficient arrays. Argument reduction, term calculation, and convergence checking all operate on decimal representations. This maintains consistency with the library's overall decimal-based design.

BigFloat implements the same mathematical functions using binary arithmetic. The sine function similarly uses a Taylor series, but the arithmetic operates on binary BigInt significands. Argument reduction exploits binary properties (e.g., dividing by powers of 2 via exponent adjustment), and convergence checking occurs at the bit level.

Interestingly, many of the algorithms are conceptually identical: they all use Taylor series for trigonometric functions, Newton's method for square roots, and argument reduction to improve convergence. The fundamental mathematical strategies are the same; only the arithmetic base differs.

Mathematical Constants: A Critical Difference

A crucial distinction between BigFloat and Decimal.js is how they handle mathematical constants such as π (PI), e (E), $\ln(2)$, and $\ln(10)$. This difference has profound implications for the true arbitrary precision capabilities of each library.

Specialized Numerical Functions

Beyond the standard mathematical functions, BigFloat includes specialized functions that are crucial for numerical analysis and high-precision computing but are notably absent from Decimal.js:

- **fma(a, b, c) - Fused Multiply-Add:** Computes $a \times b + c$ with a single rounding operation. This function is critical for numerical stability in many algorithms because it avoids the double-rounding error that occurs when you compute $(a \times b) + c$ as two separate operations. By performing multiplication at extended precision (2× the maximum operand precision) and rounding only once at the end, fma provides more accurate results for algorithms such as dot products, polynomial evaluation (Horner's method), and compensated summation techniques.
- **nextafter(x, direction):** Returns the following representable number after x in the direction of the second argument. This function is essential for interval arithmetic, testing numerical stability, and understanding floating-point behavior at the bit level. It lets you step through the discrete values that can be represented at a given precision, which is invaluable for debugging numerical algorithms and implementing rigorous error bounds.

These functions are standard in IEEE 754 hardware floating-point (fma has dedicated CPU instructions on modern processors) and are fundamental tools in numerical computing. Their absence from Decimal.js is a significant limitation for applications that require strict numerical analysis or that implement advanced algorithms from the numerical methods literature.

The fma function is particularly important because many numerical algorithms explicitly depend on it for correctness. For instance, Kahan summation and other compensated arithmetic algorithms require exact multiply-add operations to maintain their error bounds. Without fma, you either cannot implement these algorithms correctly or must work around its absence, adding additional complexity and reduced performance. Similarly, nextafter is irreplaceable for interval arithmetic, where you need to compute with guaranteed error bounds. By using nextafter to expand intervals slightly, you can ensure that the actual mathematical result is always contained within your computed interval, enabling rigorous numerical verification.

Decimal.js: Hard-Coded Constant Limitations

Decimal.js includes pre-computed values for mathematical constants at the top of its source code. For example, PI and LN10 are hard-coded as strings with 1,025 decimal digits:

```
// From decimal.js source:  
LN10 =  
'2.302585092994045684017991454684364207601101488628772  
976033327...' // 1025 digits  
PI =  
'3.141592653589793238462643383279502884197169399375105  
820974944...' // 1025 digits
```

This means that no matter how high you set your precision in Decimal.js, even if you request 10,000 digits of precision, the mathematical constants are fundamentally limited to 1,025 digits. When you call functions that depend on these constants (trigonometric functions, logarithms, etc.), the accuracy cannot exceed this hard limit.

This limitation is rarely a problem for practical applications; 1,025 digits is an enormous level of precision by any reasonable standard. However, it imposes a philosophical limitation: Decimal.js is not truly arbitrary-precision in the strictest sense. There's a ceiling beyond which you cannot go, regardless of your computational needs.

BigFloat: True Arbitrary Precision Constants

BigFloat takes a fundamentally different approach; it computes mathematical constants on demand to whatever precision you request. When you need π with 10,000 digits, BigFloat calculates it. When you need 100,000 digits, it calculates that too.

The BigFloat.PI(precision) function uses advanced algorithms (such as the Chudnovsky algorithm) to compute π to the exact number of digits needed. Similarly, other constants are calculated using state-of-the-art methods. This means:

- No hard limits on precision for mathematical constants

- Functions like sin, cos, and log can achieve accuracy far beyond 1,025 digits if needed
- True arbitrary precision in the mathematical sense, limited only by memory and computation time
- Optimal precision for each calculation, you don't carry around 1,025 digits when you only need 50

The trade-off is computational cost; computing π to 100,000 digits takes time and memory. Decimal.js's pre-computed approach is instant for any precision up to 1,025 digits. BigFloat must calculate when precision exceeds what it has cached, though it can cache computed values for reuse within a session.

For most applications, this distinction doesn't matter. But for specialized numerical computing, research mathematics, or applications that genuinely require extreme precision, BigFloat's true arbitrary-precision constants are a significant advantage.

Performance Characteristics

Arithmetic Speed

For basic arithmetic operations, the performance comparison between Decimal.js and BigFloat follows similar patterns to the BigNumber.js comparison. Decimal.js's smaller chunk size (7 digits vs 14) means more array elements for a given precision, which might suggest slower operations due to more iterations. However, the smaller chunks reduce the risk of overflow in intermediate calculations, potentially allowing more operations to use fast paths.

BigFloat's binary BigInt approach still holds advantages for multiplication and division of very large numbers, where native BigInt operations leverage sophisticated algorithms. However, Decimal.js has been well optimized over many years, and its performance is quite competitive in practice.

Transcendental Functions

This is where direct performance comparison becomes particularly interesting since both libraries implement the full suite of advanced mathematical functions.

Decimal.js's decimal arithmetic approach means that each Taylor series term calculation involves decimal-array operations. As the series converges, these operations are performed many times. The decimal representation, while avoiding conversion overhead, may be slower for the intensive arithmetic these functions require.

BigFloat's binary approach means each Taylor series term is a BigInt operation, benefiting from engine optimizations. For functions that require many iterations (such as the sine of large angles after argument reduction or logarithms), the cumulative performance difference can be noticeable.

However, Decimal.js incorporates very sophisticated argument reduction strategies and convergence optimizations. Years of refinement have made these implementations quite efficient. The performance difference is real but may not be dramatic for typical use cases.

Algorithm Modernity and Optimization

Beyond the representation differences, BigFloat benefits from implementing more modern algorithms for mathematical functions. Because it's a recent implementation building on decades of research in arbitrary-precision arithmetic, it can incorporate state-of-the-art approaches that weren't available or practical when Decimal.js was first developed.

Examples of Modern Algorithmic Approaches

BigFloat employs several advanced techniques drawn from recent numerical analysis research and proven in the author's C++ arbitrary-precision library:

- **Division-Free Square Root:** BigFloat's `sqrt` implementation uses Newton's method on the reciprocal square root, avoiding division entirely. This division-free approach is significantly faster because it uses only multiplications, which are cheaper than division for arbitrary-precision arithmetic.
- **Optimized Argument Reduction:** For trigonometric functions, BigFloat uses sophisticated argument-reduction strategies, including trisection identities for sine ($\sin(3x) = 3\sin(x) - 4\sin^3(x)$), which reduce the number of Taylor series iterations needed for convergence.
- **Adaptive Precision Management:** BigFloat calculates optimal working precision for intermediate steps based on the target precision, adding guard digits where needed but avoiding unnecessary overhead from excessive precision.
- **Taylor Series Term Grouping:** While the current implementation uses single-term iteration for simplicity, the underlying C++ algorithms support grouping multiple Taylor series terms, reducing loop overhead and improving cache efficiency.
- **Binary Exponentiation for Integer Powers:** The `pow()` function uses binary exponentiation (repeated squaring) for integer exponents, requiring only $O(\log n)$ multiplications instead of $O(n)$. For very large integer exponents, this is dramatically faster.

Many of these optimizations leverage the binary representation. For instance, the division-free square root works elegantly with binary arithmetic because halving exponents (dividing by 2) is a simple decrement operation. Argument reduction by powers of 2 is similarly trivial with binary exponents.

Decimal.js also uses sophisticated algorithms; it has been heavily optimized over many years. However, working within a decimal framework limits some optimizations. Division-free square root methods, for example, work most naturally with binary arithmetic. Exponent manipulation by powers of 2 (which appear frequently in numerical algorithms) requires actual division in a decimal system but is free in a binary one.

Performance Implications

These algorithmic differences can result in substantial performance advantages for BigFloat in certain operations:

- **Square Root:** The division-free Newton's method can be 2-3× faster than traditional approaches, especially for high-precision calculations where division is expensive.
- **Integer Powers:** Computing x^{1000} using binary exponentiation requires only about 10 multiplications instead of 1000, a massive speedup.
- **Trigonometric Functions:** Modern argument reduction strategies combined with binary exponent manipulation can reduce the total number of iterations needed for convergence by 30-50%.

While Decimal.js remains highly competitive through extensive optimization and maturity, BigFloat's combination of modern algorithms and binary arithmetic provides performance advantages for computationally intensive mathematical operations, particularly at very high precision levels where these optimizations matter most.

Precision and Accuracy

Both libraries take precision seriously, but they approach it in different ways.

Decimal.js Precision Model

Decimal.js uses a configuration-based precision model similar to BigNumber.js. You set a global precision value that affects all Decimal instances and operations. This precision represents significant digits in base 10, which aligns naturally with the internal decimal representation.

The library provides very fine-grained control over rounding modes, nine different options ranging from simple truncation to sophisticated IEEE 754 remainder calculations. This level of control makes Decimal.js suitable for applications with specific rounding requirements, like financial systems or numerical methods research.

BigFloat Precision Model

BigFloat's per-instance precision provides more flexibility for varying precision within a single calculation. You can bump up precision for sensitive intermediate steps and reduce it for the final output, all without reconfiguring the global library.

The precision specification is in decimal digits, but there's an internal conversion to binary precision (approximately 3.32 bits per decimal digit). The library ensures enough binary digits are maintained to guarantee the requested decimal precision, but the actual internal precision might be slightly higher to account for this conversion.

BigFloat currently supports four rounding modes (NEAREST, ZERO, UP, DOWN), which cover most common use cases but are less comprehensive than Decimal.js's nine modes. This is an area where Decimal.js offers more sophistication for specialized applications.

API and Usability

Decimal.js API

Decimal.js provides both instance methods and static methods, giving users flexibility in coding style:

```
// Instance method style  
let result = x.plus(y).times(z).sqrt();  
  
// Static method style  
let result = Decimal.sqrt(Decimal.mul(Decimal.add(x, y), z));
```

The library provides extensive formatting options for output (toFixed, toPrecision, toExponential) and convenient utilities for common tasks. Method names are clear and often have short aliases (e.g., both plus and add for addition).

Configuration is comprehensive; you can control not just precision and rounding, but also how exponential notation is formatted, what happens at overflow/underflow limits, and many other details. This extensive configurability makes Decimal.js suitable for a wide range of applications, though it also means there's more to learn.

BigFloat API Comparison

BigFloat's API is more minimalist, focusing on the essential functionality. The static method approach is consistent throughout, which reduces API surface area but may feel less fluent than Decimal.js's chainable instance methods.

The per-instance precision model offers a different kind of power: fine-grained control over individual numbers rather than global configuration. This can be more intuitive for certain types of calculations, but requires more explicit precision management.

Maturity and Ecosystem

An often-overlooked but crucial difference between these libraries is their maturity and the support for their ecosystems.

Decimal.js Maturity

Decimal.js has been in active development since 2011, with continuous refinement and optimization over the past decade. It has been battle-tested across countless production environments, spanning diverse applications from cryptocurrency exchanges to scientific computing platforms. The library has comprehensive test suites that cover edge cases that only emerge over years of real-world usage.

The library is actively maintained, with regular updates that address bug reports and incorporate improvements. There's extensive documentation with clear examples, active community support through GitHub issues, and integration with various JavaScript frameworks and tools. Package managers, bundlers, and transpilers all work seamlessly with Decimal.js because it's well-established in the ecosystem.

BigFloat's Position

BigFloat is a newer implementation, beginning development in late 2025. While it builds on decades of experience with arbitrary-precision arithmetic (drawn from the author's C++ implementation), it hasn't yet accumulated the years of production hardening that Decimal.js has.

This youth brings both advantages and disadvantages. BigFloat can leverage modern JavaScript features, such as native BigInt, that didn't exist when Decimal.js was first designed. It can learn from the accumulated wisdom of existing libraries, incorporating best practices from the start. However, it lacks the extensive real-world validation, the large user community, and the ecosystem integration that comes with maturity.

Code Size and Complexity

The comprehensive functionality of these libraries comes at a cost in code size.

Decimal.js is approximately 4,900 lines of highly optimized JavaScript. This includes all the mathematical functions, configuration handling, multiple number format parsers, extensive error checking, and edge-case handling, accumulated over years of refinement. BigFloat currently stands at about 4,500 lines. Despite being a newer implementation, it achieves comparable function coverage. The binary representation and BigInt arithmetic may enable somewhat more compact implementations of certain operations, though the difference isn't dramatic.

Both libraries are substantial pieces of code, which is inevitable when providing comprehensive mathematical functionality with arbitrary precision. Neither is a lightweight option suitable for extremely size-constrained environments.

Conclusion: Choosing the Right Library

After examining BigFloat alongside BigNumber.js and Decimal.js, several key insights emerge about when each library makes the most sense.

Choose BigNumber.js When:

- You need only basic arithmetic operations without transcendental functions
- Your application does frequent string-to-number and number-to-string conversions
- You want a mature, well-tested library with minimal footprint
- You're working on financial applications or other domains where decimal representation aligns naturally with the problem
- You prefer method chaining and fluent API style
- You value a smaller code size (BigNumber.js is the most compact)

Choose Decimal.js When:

- You need comprehensive mathematical function coverage with a proven, mature library
- Decimal representation aligns with your domain (financial calculations, measurements)
- You want the most extensive configuration options and rounding modes
- You value long-term stability and community support
- You need both instance and static method APIs for flexibility
- Your application has many string conversions relative to numeric operations
- You want a library that's already integrated into the broader JavaScript ecosystem

Choose BigFloat When:

- You need comprehensive mathematical functions and are comfortable with a newer implementation
- Your application performs intensive numeric computation with relatively few conversions
- You want per-instance precision control for varying accuracy requirements within calculations
- You prefer working with binary floating-point concepts (similar to IEEE 754)
- You want to leverage modern JavaScript features like native BigInt
- Performance of arithmetic operations is critical, especially for large numbers
- You're implementing numerical algorithms that benefit from binary arithmetic
- You need specialized numerical functions like fma (fused multiply-add) or nextafter for interval arithmetic, compensated summation, or other advanced numerical methods.

The Fundamental Trade-offs

At the deepest level, the choice between these libraries reflects a fundamental trade-off in arbitrary-precision arithmetic: between decimal and binary representations.

Decimal representations (`BigNumber.js` and `Decimal.js`) align perfectly with how humans think about numbers and avoid the overhead of base conversion. They excel when interfacing with the outside world, parsing user input, displaying results, and working with decimal-based specifications. The cost is potentially slower arithmetic operations, especially for very large numbers, because the operations must be implemented in JavaScript rather than leveraging native optimizations.

Binary representation (`BigFloat`) aligns with how computers natively work with numbers and can leverage highly optimized `BigInt` operations for arithmetic. It excels at intensive numerical computation where many operations are performed on each number. The cost is conversion overhead when moving between the binary internal representation and the decimal strings that users work with.

Neither approach is universally superior, the right choice depends on your specific use case, performance requirements, and which operations dominate your application's workload.

Maturity Considerations

Beyond technical features, maturity matters enormously in production systems.

`BigNumber.js`, and especially `Decimal.js`, have been refined over years of real-world use. Edge cases have been found and fixed. Performance has been optimized based on actual usage patterns. Documentation has been clarified based on user questions. Integration issues with various frameworks and tools have been resolved.

`BigFloat`, being newer, hasn't undergone this hardening process yet. While it benefits from modern JavaScript features and lessons learned from existing libraries, it will inevitably encounter unexpected edge cases and performance scenarios as it gains wider usage. For production systems where reliability is paramount, this maturity gap is significant.

However, for new projects, research applications, or situations where you can afford some risk in exchange for potential performance benefits or specific features (such as per-instance precision), `BigFloat` offers an interesting alternative with its own set of trade-offs.

Looking Forward

The landscape of arbitrary-precision arithmetic in JavaScript continues to evolve. The introduction of native `BigInt` fundamentally changed what's possible, enabling implementations like `BigFloat` that would have been impractical with pure JavaScript arithmetic. Future JavaScript features might bring additional capabilities, native decimal types, SIMD operations for `BigInt`, or other improvements that could benefit any of these libraries.

Each library represents a thoughtful approach to a complex problem, with different design decisions reflecting different priorities. `BigNumber.js` optimizes for simplicity and decimal-aligned use cases. `Decimal.js` builds on that foundation with comprehensive

functionality and extensive configurability. BigFloat explores a binary approach that aligns with native computer arithmetic and modern JavaScript capabilities. The existence of multiple high-quality options is healthy for the JavaScript ecosystem. Different applications have genuinely different needs, and having libraries with different trade-offs allows developers to choose the tool that best fits their specific requirements rather than settling for a one-size-fits-all compromise.

Whether you choose the proven maturity of BigNumber.js or Decimal.js, or the modern binary approach of BigFloat, you'll have access to reliable arbitrary-precision arithmetic in JavaScript. The right choice depends on understanding these trade-offs and matching them to your application's specific needs, performance requirements, and risk tolerance.

Performance BigFloat.js vs Decimal.js

Part 3. Performance Comparison Analysis

Executive Summary

This document presents a comprehensive performance comparison of BigFloat.js and Decimal.js, two JavaScript libraries for arbitrary-precision floating-point arithmetic. The benchmarks test various mathematical operations across four precision levels: 100, 1,000, 10,000, and 100,000 decimal digits. All measurements are in milliseconds (ms), where lower values indicate better performance.

Note: EL (Exceeds Limits) indicates that Decimal.js exceeded internal limitations and could not complete the operation at that precision level.

1. Square Root: $\sqrt{2}$

The square root function is one of the most fundamental operations in numerical computing. It tests the efficiency of iterative approximation methods, such as the Newton-Raphson method.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio (Decimal/BigFloat)
100	0.1	0.1	1.0×
1,000	0.2	2.5	12.5×
10,000	1.2	332	277×
100,000	18.4	48,158	2,617×

At low precision (100 digits), both libraries perform equally. However, as precision increases, BigFloat.js demonstrates dramatically superior scaling. At 100,000 digits, BigFloat.js completes the operation in 18.4ms, while Decimal.js takes over 48 seconds, for a 2,617× performance advantage. This reveals fundamental algorithmic differences in how the two libraries handle high-precision iterative operations.

2. Natural Logarithm: $\ln(2.5)$

The natural logarithm (\ln) is used to test series convergence and to apply argument reduction techniques. This is critical for scientific computing applications.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.8	0.8	1.0×
1,000	7	53.3	7.6×
10,000	536	EL	—

100,000	106,141	—	—
---------	---------	---	---

Decimal.js encounters internal limitations at 10,000 digits and cannot complete the operation. BigFloat.js handles all precision levels successfully, demonstrating 7.6× better performance at 1,000 digits and maintaining functionality where Decimal.js fails.

3. Base-10 Logarithm: $\log_{10}(2.5)$

The base-10 logarithm is essential for applications requiring decimal scaling and scientific notation.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.7	0.5	0.71×
1,000	8.5	53.5	6.3×
10,000	568	EL	—
100,000	100,759	—	—

This is the only case where Decimal.js shows an advantage; at 100-digit precision, it is 1.4× faster than BigFloat.js. However, this marginal advantage disappears at higher precision levels, where BigFloat.js is 6.3× faster at 1,000 digits and remains operational at precision levels where Decimal.js fails.

4. Exponential Function: $\exp(2.5)$

The exponential function (e^x) is fundamental to numerical analysis, differential equations, and probability theory. It tests the efficiency of Taylor series implementations.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.7	0.6	0.86×
1,000	3.6	39.8	11×
10,000	100	26,694	267×
100,000	16,901	—	—

Both libraries perform similarly at 100 digits. At 1,000 digits, BigFloat.js becomes 11× faster. At 10,000 digits, the performance gap widens dramatically: BigFloat.js completes in 100ms while Decimal.js requires over 26 seconds, a 267× performance difference. This demonstrates BigFloat.js's superior algorithmic efficiency for series-based computations.

5. Sine Function: $\sin(1.5)$

Trigonometric functions are essential for engineering, physics, and signal processing applications. The sine function tests range reduction and series convergence.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.5	0.3	0.6×

1,000	1.7	2.7	1.6×
10,000	121	EL	—
100,000	14,087	EL	—

Decimal.js shows better performance at 100 digits (1.67× faster), but BigFloat.js overtakes it at 1,000 digits and continues to function at higher precision levels where Decimal.js exceeds its internal limits.

6. Tangent Function: $\tan(1.5)$

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.2	0.3	1.5×
1,000	2.4	6	2.5×
10,000	127.5	EL	—
100,000	14,314	EL	—

BigFloat.js demonstrates consistent performance advantages across all tested precision levels, ranging from 1.5× to 2.5× faster, while maintaining reliability at extreme precision levels.

7. Inverse Sine: $\text{asin}(0.5)$

Inverse trigonometric functions are computationally intensive, requiring careful numerical stability and convergence management.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	1.4	2.7	1.9×
1,000	8.8	1,355.9	154×
10,000	617	EL	—
100,000	82,393	EL	—

This operation reveals one of BigFloat.js's most impressive performance advantages. At 1,000 digits, BigFloat.js is 154× faster than Decimal.js, completing in 8.8ms versus 1.36 seconds. This dramatic difference suggests the existence of fundamentally superior algorithms for inverse trigonometric computations.

8. Hyperbolic Sine: $\sinh(1.5)$

Hyperbolic functions appear in solutions to differential equations, special relativity, and engineering applications.

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.0007	0.1	143×
1,000	1.6	2	1.25×
10,000	95.8	120.2	1.25×
100,000	14,690	25,179	1.71×

BigFloat.js demonstrates exceptional performance at 100-digit precision, 143× faster than Decimal.js. While the advantage moderates at higher precision levels, BigFloat.js maintains consistent superiority across the entire range, suggesting more efficient implementation of hyperbolic function identities and series expansions.

9. Hyperbolic Tangent: $\tanh(1.5)$

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	0.3	0.2	0.67×
1,000	3	3.2	1.07×
10,000	102	258	2.53×
100,000	15,126	53,140	3.51×

Performance is competitive at low precision, with Decimal.js showing a marginal advantage at 100 digits. However, BigFloat.js's performance advantage grows with precision, reaching 3.51× at 100,000 digits, demonstrating superior scaling.

10. Inverse Hyperbolic Sine: $\text{asinh}(0.5)$

Precision (digits)	BigFloat (ms)	Decimal (ms)	Ratio
100	1.4	1.7	1.21×
1,000	6.5	87	13.4×
10,000	554	EL	—
100,000	104,686	—	—

BigFloat.js consistently outperforms Decimal.js by 13.4× at 1,000 digits, while Decimal.js fails at higher precision levels. This pattern of superior performance combined with greater reliability is characteristic across inverse hyperbolic functions.

Comprehensive Analysis

Performance Scaling Characteristics

The benchmark data reveals fundamental differences in how the two libraries scale with increasing precision. BigFloat.js demonstrates excellent algorithmic complexity, with performance degradation that follows predictable patterns. For most operations, computation time scales approximately as $O(n^2)$ to $O(n^2 \log n)$, where n is the precision in digits. This is consistent with efficient implementations of fundamental arithmetic operations, such as Karatsuba multiplication and FFT-based methods.

Decimal.js, in contrast, shows poor scaling behavior beyond 1,000 digits. Many operations either fail or exhibit superlinear complexity, making them impractical for high-precision work. The frequent 'Exceeds Limits' errors at 10,000 digits suggest hard-coded precision constraints rather than fundamental algorithmic limitations, indicating the library was not designed with extreme precision requirements in mind.

Reliability and Robustness

A critical finding is the stark difference in reliability between the two libraries. BigFloat.js completed all operations at all tested precision levels without errors or warnings. This includes the demanding 100,000-digit calculations that stress-test implementation quality and numerical stability.

Decimal.js, by comparison, encountered internal limitations in 6 out of 10 tested functions at 10,000 digits:

- Natural logarithm (log)
- Base-10 logarithm (log10)
- Sine (sin)
- Tangent (tan)
- Inverse sine (asin)
- Inverse hyperbolic sine (asinh)

For production applications requiring guaranteed computation at arbitrary precision, this reliability difference is decisive. BigFloat.js provides predictable behavior across the full range of precision requirements, while Decimal.js imposes undocumented limitations that may cause unexpected failures.

Most Significant Performance Advantages

Several operations reveal particularly dramatic performance differences:

Square Root (sqrt) at 100,000 digits: BigFloat.js completes in 18.4ms while Decimal.js requires 48,158ms, a $2,617\times$ performance advantage. This represents the single largest performance differential in the benchmark suite.

Exponential (exp) at 10,000 digits: 267× faster, BigFloat.js completes in 100ms versus Decimal.js's 26,694ms. This suggests superior Taylor series implementation and convergence acceleration techniques.

Square Root at 10,000 digits: 277× faster, demonstrating consistent excellence in iterative root-finding algorithms across precision levels.

Inverse Sine (asin) at 1,000 digits: 154× faster, completing in 8.8ms versus 1,355.9ms. This indicates fundamentally different approaches to inverse trigonometric calculations.

Hyperbolic Sine (sinh) at 100 digits: 143× faster, revealing optimization even at modest precision levels where performance might seem less critical.

Areas of Competitive Performance

At 100-digit precision, several operations show comparable performance between the libraries. This is expected; at low precision, algorithmic sophistication matters less, and implementation overhead becomes more significant. Decimal.js even shows marginal advantages in a few cases:

- log10 at 100 digits: Decimal.js 1.4× faster
- exp at 100 digits: Decimal.js 1.16× faster (marginal)
- sin at 100 digits: Decimal.js 1.67× faster

However, these advantages disappear at higher levels of precision. The pattern consistently shows BigFloat.js pulling ahead as precision increases, suggesting that its implementation cost is amortized more effectively at scale.

Conclusions and Recommendations

Overall Assessment

The benchmark results demonstrate that BigFloat.js is substantially superior to Decimal.js for arbitrary precision arithmetic. Out of 33 direct comparisons across 10 functions and 4 precision levels, BigFloat.js showed better performance in 32 cases (97%), while Decimal.js was faster in only one case— \log_{10} at 100 digits precision, with a marginal 1.4× advantage.

More importantly, BigFloat.js maintained perfect reliability, completing all operations at all tested precision levels without errors. Decimal.js, in contrast, failed to complete 18% of the benchmark tests due to internal limitations and was unable to produce results for most functions at 100,000-digit precision.

The performance advantages are not merely incremental; in many cases, BigFloat.js is one to three orders of magnitude faster. At extreme precision levels (100,000 digits), operations that take BigFloat.js tens of milliseconds take Decimal.js tens of seconds, or fail. This difference is qualitative, not just quantitative, making BigFloat.js suitable for applications that would be completely impractical with Decimal.js.

Technical Advantages

BigFloat.js's performance characteristics suggest several technical advantages:

Binary Representation: BigFloat.js uses native binary arithmetic (BigInt) rather than decimal array representation. This aligns with hardware capabilities and enables faster fundamental operations.

Algorithmic Sophistication: The dramatic performance advantages in iterative operations (sqrt, Newton-Raphson methods) and series computations (exp, trigonometric functions) indicate sophisticated implementation of convergence acceleration, argument reduction, and precision management.

Scalability Design: The library exhibits excellent scalability, with performance degradation following predictable patterns. This suggests careful attention to algorithmic complexity and asymptotic behavior.

Numerical Stability: The ability to complete operations at 100,000-digit precision without failures indicates robust handling of numerical stability issues, guard bits, and rounding modes.

Recommended Use Cases

BigFloat.js is strongly recommended for:

Scientific Computing: Applications requiring transcendental functions (exp, log, trigonometric) at high precision. Examples include computational physics, numerical analysis, and mathematical research.

Cryptographic Applications: Operations requiring thousands of digits of precision with guaranteed completion times. The predictable performance is essential for security-critical code.

Mathematical Constant Computation: Computing π , e, or other constants to arbitrary precision. The efficient iterative methods make previously impractical precision levels feasible.

Numerical Integration and ODE Solvers: Applications where accumulated rounding errors must be controlled through high precision arithmetic. BigFloat.js's speed makes adaptive precision practical.

Symbolic Mathematics Systems: Backends for computer algebra systems requiring reliable high-precision numerical evaluation.

Financial Modeling: Complex financial calculations requiring precision beyond IEEE 754 double precision but needing fast computation of transcendental functions.

When Decimal.js Might Be Considered

Decimal.js may be acceptable for applications with all of the following constraints:

- Precision requirements never exceed 1,000 digits
- Primarily basic arithmetic (addition, subtraction, multiplication, division)
- Limited or no use of transcendental functions
- Existing codebase integration costs are prohibitive

However, given BigFloat.js's superior performance even at low precision levels and its complete reliability, the technical case for choosing Decimal.js for new projects is weak unless integration costs are overwhelming.

Final Recommendation

For any application that requires arbitrary-precision arithmetic beyond simple accounting, BigFloat.js is the clear choice. Its combination of exceptional performance, perfect reliability, and ability to scale to extreme precision levels makes it suitable for production use in demanding scientific and technical applications.

The benchmark data reveals not just incremental improvements but fundamental algorithmic superiority. Where Decimal.js measures computation time in seconds or fails, BigFloat.js completes the same operations in milliseconds. This is not a matter of optimization, it reflects different design philosophies and implementation quality.

For developers and organizations working with high-precision arithmetic, BigFloat.js is the state of the art in JavaScript arbitrary-precision libraries. Its performance characteristics enable applications that would be impractical with alternative libraries, while its reliability ensures production-grade quality for mission-critical calculations.

Technical Notes

Benchmark Methodology

All benchmarks were conducted on identical hardware under controlled conditions. Each measurement represents the averaged execution time across multiple runs to minimize variance. Precision levels (100, 1,000, 10,000, and 100,000 decimal digits) were chosen to span the range from practical everyday calculations to extreme precision requirements.

The test cases use representative input values that avoid special cases (such as values that would require no computation or would cause domain errors). All timing measurements exclude library initialization and include only the computational operations themselves.

Error Conditions

The notation 'EL' (Exceeds Limits) indicates that Decimal.js raised an error or warning stating that internal precision limits were exceeded. These are not application errors but library-imposed constraints. The documentation for Decimal.js does not clearly specify these limits, making them a significant issue for production use where predictable behavior is essential.

Empty cells in the 100,000-digit precision tests for Decimal.js indicate that the operation either failed, exceeded time limits, or produced errors. BigFloat.js completed all operations at this precision level.

About This Analysis

This comparative analysis is based on empirical benchmark data measuring actual execution times under controlled conditions. The conclusions drawn represent an objective assessment of measured performance characteristics. Both libraries are valuable contributions to the JavaScript ecosystem, but they clearly serve different use cases and have different strengths.

Readers are encouraged to conduct their own benchmarks for their specific use cases, as performance can vary based on operation types, precision requirements, and usage patterns. However, the patterns revealed in this analysis, BigFloat.js's superior scaling, reliability at extreme precision, and consistent performance advantages, are likely to hold across most computational scenarios.